



DEPARTMENT OF ELECTRICAL AND
ELECTRONIC ENGINEERING

FINAL YEAR PROJECT

**A Data-Movement Characterization of a
Production Visual-SLAM Workload for
Memory-Centric Heterogeneous Architectures**

21/ENG/058 – K.R.M.I. Jayasiri

21/ENG/066 – K.A.K.T. Gihan

21/ENG/096 – P.R.C.J. Bandara

Supervised by

Dr. U.L. Wijewardhana

Acknowledgments

We thank our supervisor, Dr. U.L. Wijewardhana, for his guidance and feedback throughout this project, and for the latitude to pursue a measurement-first approach to the problem.

We are grateful to the Department of Electrical and Electronic Engineering, University of Sri Jayewardenepura, for the facilities and support that made the experimental work in this thesis possible.

This thesis builds on NVIDIA’s public release of cuVSLAM and the Isaac robotics stack, and on the KITTI, EuRoC MAV, TUM RGB-D/VI, and ICL-NUIM benchmark datasets, whose creators made their data freely available to the research community. We also thank the developers of Nsight Compute, Nsight Systems, and NVBit, the instrumentation this characterization is built on.

Finally, we thank our families and friends for their patience and support throughout this project.

Contents

1	Introduction	12
1.1	Motivation: Physical AI Faces the Memory Wall	12
1.2	Problem Statement	13
1.3	Research Questions	14
1.4	Contributions	14
1.5	Scope and Non-Goals	16
1.6	Thesis Organization	16
2	Literature Review	17
2.1	Background	17
2.1.1	The GPU Execution Model	17
2.1.2	The Memory Hierarchy and the Five Address Spaces	18
2.1.3	The Memory Wall	19
2.1.4	Visual SLAM and cuVSLAM	20
2.1.5	Candidate Substrates for Memory-Centric Execution	21
2.1.6	The Measurement Instruments	22
2.2	Related Work	23
2.2.1	Data-Movement Characterization Methodologies	23
2.2.2	Processing-in-Memory and Near-Data Systems	24
2.2.3	SLAM Systems and Their Evaluation	25
2.2.4	GPU Instrumentation	25
2.2.5	Positioning	25
3	Methodology	27
3.1	The Six-Step Pipeline	27
3.2	Experimental Setup	28
3.2.1	Hardware	28
3.2.2	Locked Clocks and Measured Ceilings	28
3.2.3	Workloads and Datasets	29
3.2.4	Capture Discipline	29
3.3	Metrics	30
3.3.1	Speed-of-Light Utilizations	30

3.3.2	LFMR (GPU Form)	30
3.3.3	Memory Intensity (MPKI)	30
3.3.4	Arithmetic Intensity and the Roofline	30
3.3.5	Occupancy and the Stall Taxonomy	31
3.3.6	Coalescing	31
3.3.7	Locality (Trace-Side)	31
3.4	The Classifier and Taxonomy	32
3.4.1	Thresholds	32
3.4.2	The Ordered Rules	32
3.4.3	Sensitivity Analysis	34
3.5	From Class to Verdict: Substrate and Fault Mapping	34
3.6	The Placement Model	35
3.7	DAMOV Parity: The Completeness Argument	35
4	Results	37
4.1	Measurement Validity: Accuracy, Neutrality, and Noise	37
4.1.1	Accuracy Validation: The 141-Configuration Matrix	37
4.1.1.1	Discrepancies and Their Causes	37
4.1.2	Profiling is Accuracy-Neutral	38
4.1.3	Nondeterminism, Characterized	39
4.1.4	Summary of the Trust Chain	40
4.2	Kernel-Level Characterization	40
4.2.1	What the Screen Reveals	40
4.2.2	Class Stability Across Inputs	41
4.2.3	The Taxonomy Is Discovered, Not Asserted	41
4.2.4	Cross-Microarchitecture Agreement	42
4.2.5	The Calibration Suite: Classifying Known Truth	42
4.2.6	Near-Data Affinity of the Workload	43
4.2.7	The System Edge: Host–Device Transfers	43
4.3	Machine-Independent Locality Analysis	44
4.3.1	Trace-Side Metrics	44
4.3.2	The Front-End Is Cache-Immune Streaming	45
4.3.3	The Loop-Closure Scan: A Self-Correction on Record	46

4.3.4	Session-Scale Growth and Migration	47
4.4	Data-Structure Attribution: Naming Every Byte	48
4.4.1	The Three-Layer Instrument	48
4.4.2	Coverage: Auditing the Windows	49
4.4.3	The GPU Memory Budget Is Static	49
4.4.4	Attribution Results	49
4.4.5	The Register-Spill Finding	51
4.4.6	Bounded Residuals	51
4.5	Synthesis: Substrate Verdicts and Measured Baselines	52
4.5.1	The Three-Way Split	52
4.5.2	Verdict Dynamics Across Workloads	52
4.5.3	The Analytical Placement Bound	53
4.5.4	The Measured Energy Baseline	54
4.5.5	The Host-Side Boundary	55
4.5.6	Faults Named for the Kernels That Stay	55
5	Discussion	56
5.1	Internal Validity	56
5.2	External Validity	57
5.3	Construct Validity	57
5.4	What Is Out of Scope, and Why	58
6	Conclusions	59
6.1	What Was Established	59
6.2	The Thesis, Restated	60
7	Recommendations and Future Work	61
7.1	Future Work	61
8	References	62
9	Appendices	66
9.1	Glossary	66
9.2	Artifact and Reproducibility Description	68

9.2.1	Layout	69
9.2.2	Environment	70
9.2.3	Reproduction paths	70
9.3	Extended Result Tables	71
9.3.1	Cross-sequence attribution consistency (excerpt)	71
9.3.2	Instrumented vs. baseline trajectories (QoR excerpt)	71
9.3.3	KITTI per-sequence odometry accuracy	71
9.3.4	Measured stage map (NVTX, excerpt)	71

List of Figures

1	The GPU memory hierarchy.	19
2	The GPU-DAMOV six-step pipeline.	27
3	Measured roofline of cuVSLAM kernels.	31
4	The ordered classification tree (G0–G7).	33
5	Kernel distribution across the taxonomy.	42
6	Reuse-distance CDF: cache-immune streaming.	45
7	Per-kernel memory-space attribution.	51
8	The three-way substrate split.	53

List of Tables

2	Classifier thresholds (values as frozen in <code>analysis/classify.py</code>). . .	32
3	The GPU bottleneck taxonomy and its default near-data affinity. . .	34
4	Substrate verdict mapping.	35
5	Converged mode-average accuracy vs. the cuVSLAM technical report (RMSE APE in metres, odometry/SLAM). The characteri- zation uses the stereo and RGB-D modes.	38
6	The <code>st_track_with_cache</code> correction: blended-space analysis vs. space-filtered re-derivation (TUM room-scale / KITTI street-scale). . .	47
7	The static GPU allocation budget (full session; peak = total, nothing freed mid-run).	50
8	The three-way substrate split, with measurement backing each claim.	54
9	Attribution consistency across 27 sequences (excerpt).	71
10	Quality-of-results check: instrumented (TaggedAllocator) build vs. baseline build, same configurations (APE, m).	72
11	KITTI stereo odometry, per sequence: absolute error and segment- relative drift (100–800 m KITTI protocol).	72
12	Kernel-to-stage mapping, measured from cuVSLAM’s own NVTX ranges.	73

List of Abbreviations

AI	Arithmetic Intensity (FLOPs per DRAM byte)
APE / ATE	Absolute Pose / Trajectory Error
ARI	Adjusted Rand Index
BA	Bundle Adjustment
CDF	Cumulative Distribution Function
CoV	Coefficient of Variation (std. dev. $\div$ mean)
CUDA	Compute Unified Device Architecture (NVIDIA GPU programming model)
DAMOV	the Data-Movement bottleneck methodology of Oliveira et al. [1]
DRAM	Dynamic Random-Access Memory (GPU main/“global” memory)
FLOP	Floating-Point Operation
FMA	Fused Multiply–Add
GFTT	Good Features To Track (Shi–Tomasi corner detector)
GPU	Graphics Processing Unit
H2D / D2H	Host-to-Device / Device-to-Host (PCIe copies)
IMU	Inertial Measurement Unit
ISP	In-Storage Processing (computational storage); in camera contexts, Image Signal Processor
LFMR	Last-to-First Miss Ratio
LK	Lucas–Kanade optical flow
LMDB	Lightning Memory-mapped Database
MPKI	(DRAM sector) Misses Per Kilo-Instruction
NDP	Near-Data Processing
ncu	NVIDIA Nsight Compute (hardware-counter profiler)
nsys	NVIDIA Nsight Systems (timeline profiler)
NVBit	NVIDIA Binary Instrumentation Tool
NVTX	NVIDIA Tools Extension (code-range annotations)
PCIe	Peripheral Component Interconnect Express (CPU–GPU link)

PiM / PnM	Processing-in-Memory / Processing-near-Memory
QoR	Quality of Results (instrumented-vs-baseline output equivalence)
RMSE	Root-Mean-Square Error
RPE / RTE	Relative Pose / Trajectory Error
SLAM	Simultaneous Localization And Mapping
SM	Streaming Multiprocessor
SoL	Speed-of-Light (utilization as % of theoretical peak)
SRAM	Static Random-Access Memory (on-chip storage)
VIO	Visual-Inertial Odometry
VRAM	Video RAM (the GPU's DRAM)

Abstract

Physical AI systems such as robots, drones, and autonomous vehicles run their perception stacks on GPUs. For these systems, the dominant cost is increasingly not arithmetic, but data movement: the time and energy spent transferring bytes between memory and compute resources. Memory-centric hardware including processing-in-memory (PiM), in/near-storage processing, and near-sensor computation aims to address this cost. However, hardware cannot be optimally designed for workloads that have not been measured.

This thesis presents the first per-kernel, per-data-structure memory characterization of a *production* visual-SLAM system: cuVSLAM, the CUDA-accelerated simultaneous localization and mapping library deployed in NVIDIA’s Isaac robotics stack.

The thesis makes three primary contributions:

1. **Methodology:** It introduces GPU-DAMOV, a re-derivation of the DAMOV data-movement methodology for GPUs. In this approach, the last-to-first miss ratio is redefined for a two-level cache hierarchy. Latency-boundedness is conditioned on occupancy, and memory-divergence (coalescing) is introduced as a first-class axis with no CPU analogue.
2. **Empirical Analysis:** The method is applied to cuVSLAM across 27 sequences from four public datasets (KITTI, EuRoC, TUM RGB-D, TUM-VI), with zero failed runs. This yields an eight-class bottleneck taxonomy (G0–G7) that is stable across inputs (91% cross-sequence class consistency), independently validated by two unsupervised clustering algorithms (k-means and Ward hierarchical; purity ≈ 0.68), and 80% reproducible across two GPU microarchitectures.
3. **Trust and Validation:** A novel allocation-journaling instrumentation (TaggedAllocator), combined with binary-instrumentation address traces, attributes every byte of DRAM traffic to a named data structure. All 274 allocations are resolved, and 48 of 49 kernels carry a unanimous data-structure tag across all sequences.

Additionally, the thesis demonstrates that the profiled system’s trajectory accuracy reproduces NVIDIA’s published results across a 141-configuration matrix. Profiling itself is shown to be accuracy-neutral on deterministic modes: the profiled trajectories are bit-identical to unprofiled runs across a 192-configuration campaign.

The characterization resolves cuVSLAM’s memory behaviour into a three-way split with direct architectural consequences. The image front-end is *cache-immune streaming*: its measured reuse-distance profile is flat from 64 KiB to 48 MiB, so no realizable cache helps, and the measured 1.68 MB-per-frame sensor upload argues for near-sensor consumption. The bundle-adjustment solver is hot-persistent and largely compute-efficient: its DRAM traffic is 97% one named, pre-sized linear-system structure. The loop-closure back-end is a scattered gather (23–30 sectors per warp access) over a keyframe database whose GPU-resident state is a fixed 6.7 MB buffer while the session-scale store grows host-side (the in-storage-processing case); and 92.6% of the scan kernel’s DRAM volume is register-spill traffic, a compiler artifact invisible to counter-only studies. In total, 60–72% of GPU time carries PiM affinity under a sensitivity-tested classification. A measured whole-run energy baseline (34.7 J) and host-side I/O characterization (66 MB storage ingestion; 708 MB peak host memory) complete the evidence base for the follow-on architecture study, for which simulator-configuration groundwork is laid.

Beyond its findings, the thesis documents two self-corrections in which independent measurement methods initially disagreed and were reconciled (one reversing a claim made earlier in the project), and argues that this disagree-then-agree arc, with all artifacts on record, is precisely what distinguishes measurement science from advocacy.

1 Introduction

1.1 Motivation: Physical AI Faces the Memory Wall

Every robot navigating its surroundings requires instant answers to the two core questions: *Where am I?* And *What does my surroundings look like?* Visual SLAM is the computational domain tasked with answering both using a video stream from cameras. In commercial robotics, these visual SLAM operations are usually performed on GPUs. This thesis studies cuVSLAM [2]: a GPU-accelerated visual SLAM library running in the Isaac robotics stack on a physical robot. By definition, this type of workload falls into the category of so-called *Physical AI*, systems which are capable of interacting with the real world and complying with latency and energy constraints.

For this type of robot, the major constraint on processors is not arithmetic. With generations of hardware advancing, floating point calculations have become virtually costless compared to the huge amount of energy needed to read and write them from off-chip DRAM: a memory read operation consumes orders of magnitude more energy than computations do. This problem, called the *memory wall* [3], is the modern incarnation of the von Neumann bottleneck. For battery-operated robots, the implication is direct: joules spent reading and writing pixels and map entries are joules lost to either flight time or actuation.

The architecture community’s response has been the development of *memory-centric* hardware. This includes processing-in-memory (PiM), which places compute engines at or inside DRAM devices [4], [5]; in-storage or near-storage processing, which executes queries where the data is stored [6], [7]; and near-sensor processing, which consumes pixel streams before they reach main memory. Commercial implementations exist for each of these approaches [8], [9], [10]. However, there is still no principled method to determine, for a given production workload, *which* computations belong on *which* substrate. Answering this question requires precise measurement: a *characterization* of exactly how the workload moves data.

1.2 Problem Statement

The problem addressed in this thesis can be summarized in a single sentence:

For each GPU kernel of a production visual-SLAM system, determine by measurement, at the granularity of named data structures, and with quantified confidence whether its bottleneck is data movement, which class of data movement, and, consequently, which hardware substrate (GPU, CPU, DRAM-PiM, scatter-capable PiM, near-sensor, or in/near-storage) is the appropriate home for it.

Three properties make this problem challenging, and each motivates a key contribution of this thesis:

1. **The reference methodology is CPU-centric.** The established data-movement characterization methodology, DAMOV [1], was developed for CPUs. GPUs violate its assumptions in specific ways: they have a two-level (not three-level) cache hierarchy; they hide memory latency with massive thread parallelism rather than stalling; and they exhibit a failure mode with no CPU analogue *memory divergence*, where the 32 threads of a warp scatter their accesses, wasting up to $8\times$ the fetched bandwidth. A GPU characterization must adapt the methodology, not merely reuse it (see Chapter 3).
2. **Kernel-level claims are not architect-level claims.** While a profiler can report that a kernel is memory-bound, an architect needs to know *which data structure* the traffic is associated with such as an image, a solver matrix, or a growing keyframe database because the appropriate substrate depends on this distinction. Attributing every byte of DRAM traffic to a named allocation in a closed-source production binary requires purpose-built instrumentation (see Chapter 4.4).
3. **Measurements must be proven trustworthy.** A critical reviewer will ask: Was the workload computing correctly? Did the instrumentation alter system behavior? Are the measurements stable across runs? Would the classification hold under different thresholds, clustering algorithms, or hardware? Each of these concerns is addressed with dedicated experiments in

this thesis (see Chapters 4.1 and 4.2).

1.3 Research Questions

- RQ1:** Can DAMOV’s data-movement methodology be re-derived for GPUs such that each of its components screening, machine independent locality analysis, bottleneck classification, and robustness validation has a measured GPU analogue? (See Chapter 3; validated in Chapters 4.2–4.3.)
- RQ2:** What are the memory-bottleneck classes of a production visual-SLAM workload, and are these properties intrinsic to the kernels or artifacts of particular inputs, thresholds, or hardware? (See Chapter 4.2.)
- RQ3:** Which named data structures receive each kernel’s DRAM traffic, and what fraction of that traffic is program data versus compiler-generated scratch? (See Chapter 4.4.)
- RQ4:** What is the resulting substrate placement for the workload’s computation? What fraction of GPU time is offload-eligible, and what measured baselines (energy, host I/O) must a follow-on architecture study improve upon? (See Chapter 4.5.)
- RQ5:** Can the entire measurement apparatus be shown to be accuracy-neutral, that is, can it be demonstrated that the characterized system functions correctly, with outputs unchanged by observation? (See Chapter 4.1.)

1.4 Contributions

1. **GPU-DAMOV, a new methodology** (Chapter 3). This is a component-by-component re-derivation of DAMOV for GPUs: LFMR is redefined for a two-level hierarchy; latency-boundedness is conditioned on occupancy; a coalescing axis is added; and DAMOV’s simulated intervention experiment is replaced by two *measured* interventions (a clock-domain sensitivity sweep and a direct reuse-distance cache-capacity curve), with a simulated third axis prepared as a generated configuration. The method’s endpoint extends beyond DAMOV: not merely “memory-bound,” but a per-kernel *substrate*

verdict plus, when the verdict is “stay on GPU,” a named architectural fault to address.

2. **The first per-data-structure memory characterization of a production SLAM system** (Chapters 4.2–4.5). The study covers 49 kernels $\times$ 27 sequences $\times$ 4 datasets $\times$ 192 configuration variants, producing an eight-class taxonomy (G0–G7) where class membership is 91% consistent across sequences, independently recovered by two clustering algorithms, and 80% reproducible across two GPU microarchitectures.
3. **TaggedAllocator attribution** (Chapter 4.4). This three-layer instrumentation source-level allocation journaling with call stacks, driver-level allocation lifetime capture, and a streaming join against binary-instrumentation address traces resolves all 274 allocations and gives 48 of 49 kernels a unanimous data-structure tag. It reveals a finding invisible to counters: 92.6% of the loop-closure scan’s DRAM volume is register-spill scratch, not data.
4. **A validated trust chain** (Chapter 4.1). The thesis reproduces the vendor’s published trajectory accuracy with a 141-configuration accuracy matrix; demonstrates profiling is accuracy-neutral (bit-identical trajectories on deterministic modes) in a 192-configuration campaign; measures run-to-run noise (CoV 0.14% at locked clocks); and includes a $\pm 25\%$ threshold sensitivity analysis.
5. **A reusable, workload-agnostic toolchain** (released under the MIT license). This includes a single-file experiment runner, a configuration-mutation regime, a three-profiler harness with adapters for arbitrary GPU workloads, a stdlib-only analysis library that regenerates every table from committed CSVs without a GPU, and an interactive evidence dashboard.
6. **Two documented self-corrections** in which independent methods initially disagreed but were reconciled one reversing an earlier project-internal conclusion (see Chapter 4.3) are presented as a feature of the methodology: the pipeline is designed to catch its own errors.

1.5 Scope and Non-Goals

This thesis is a *characterization*, deliberately stopping at measured candidacy and an analytical placement bound. It does *not* simulate a PiM or in-storage design and, therefore, reports no end-to-end speedup or energy delta for the proposed substrates; such work is reserved for the follow-on architecture study, for which this thesis generates simulator configurations and measures baselines (see Section 7.1). The focus is on one workload family cuVSLAM chosen for its representativeness in production deployment rather than breadth across all SLAM implementations; the toolchain’s adapter mechanism provides a path to broader applicability. Hardware evaluation uses one desktop-class GPU as the primary instrument, with a second microarchitecture used to validate classification across devices.

1.6 Thesis Organization

Chapter 2.1 builds the technical background from first principles: the GPU execution model, the memory hierarchy and its five address spaces, SLAM and cuVSLAM, candidate substrates, and measurement instruments. Chapter 2.2 situates the work among data-movement characterizations, PiM systems, and SLAM acceleration studies. Chapter 3 presents GPU-DAMOV: metrics, decision tree, taxonomy, placement model, and experimental setup. Chapter 4.1 establishes trust: accuracy validation, profiler neutrality, and measurement rigor. Chapter 4.2 reports the kernel-level characterization and its four-way validation. Chapter 4.3 details the machine-independent locality analysis and the first self-correction. Chapter 4.4 presents the data-structure attribution and register-spill finding. Chapter 4.5 synthesizes the substrate verdicts, placement model, and measured energy and host-I/O baselines. Chapter 5 enumerates threats to validity. Chapter 6 concludes with a phased roadmap. Appendices provide a glossary, artifact/reproducibility description, and extended result tables.

2 Literature Review

2.1 Background

This chapter develops every concept the thesis relies on, in logical order: the GPU execution model (§2.1.1), the memory hierarchy and its address spaces (§2.1.2), why data movement dominates cost (§2.1.3), visual SLAM and cuVSLAM (§2.1.4), the candidate hardware substrates (§2.1.5), and the measurement instruments (§2.1.6).

2.1.1 The GPU Execution Model

Whereas a CPU is designed to minimize latency of an instruction stream, a GPU is built with the aim of maximizing throughput of tens of thousands of streams simultaneously. CUDA is used to program NVIDIA GPUs: the developer writes a *kernel*, an arbitrary single function, which is launched in a grid with thousands to millions of *threads* handling different pieces of data. These threads are grouped into statically defined clusters of 32 threads each called *warps*. The 32 *lanes* inside each warp compute the same instruction over different data. The warps are arranged into thread blocks that, in turn, are mapped onto *streaming multiprocessors* (SMs), fundamental compute units of the GPU. The hardware under test in this research is an NVIDIA RTX 2000 Ada generation GPU (compute capability `sm_89`) with 22 SMs, 25 MB L2 shared cache, and 16 GB GDDR6 DRAM.

Two warp-level behaviors are central throughout this thesis:

Divergence. If a branch disables some lanes of a warp, the disabled lanes perform no useful work. The counter `mean_active_lanes` measures this; 32.0 means full convergence. Every cuVSLAM kernel measured in this thesis operates at 32.0 active lanes divergence is *not* an issue, immediately shifting focus to memory behavior.

Coalescing. When a warp issues a load, the hardware merges the 32 lane addresses into fetches of 32-byte *sectors* (four sectors form a 128-byte cache line). If the lanes address 32 consecutive words, a few sectors satisfy the entire warp a *coalesced* access. If the lanes are scattered, the hardware may fetch up to 32

sectors and discard most of each one a *scattered gather* that wastes up to $8\times$ the useful bandwidth. The metric *sectors per request* quantifies this (≈ 4 for perfectly coalesced 4-byte accesses; $\rightarrow 32$ for fully scattered). Coalescing has no CPU analogue and is one of the three reasons DAMOV cannot be directly applied to GPUs (see Chapter 3).

The term *occupancy* defines the ratio of the number of resident warps to the maximum number of concurrent warps that an SM can hold. This is the GPU’s primary way of latency masking: if one warp gets stalled because of memory access, another one is executed. While a high occupancy level can completely hide the memory wall *latency*, it cannot make up for the memory *bandwidth* limit. This inherent tradeoff, while latency can be masked, bandwidth cannot, is the core of the bottleneck classification presented in Chapter 3.

2.1.2 The Memory Hierarchy and the Five Address Spaces

GPU storage forms a hierarchical pyramid (Figure 1): per-thread *registers* (the fastest and smallest); per-SM on-chip storage of 100 KB, serving as both L1 cache and programmer-managed *shared memory*; the chip-wide L2 cache (25 MB); the GPU’s DRAM (16 GB, providing a measured 205 GB/s on the locked instrument, with latency roughly $100\times$ that of a register); host RAM via the PCIe bus; and persistent storage. Each successive level increases the energy cost per byte by an order of magnitude.

Beyond this physical pyramid, every memory *instruction* addresses one of several logical

spaces. Confusing these spaces can invalidate analyses (see Chapter 4.3 for a demonstration):

Global The program’s actual data in DRAM everything allocated with `cudaMalloc`.
Instructions: LDG/STG. All DRAM traffic and locality claims in this thesis pertain to global space.

Shared The on-chip per-SM scratchpad (LDS/STS), used for convolution tiles, sort staging, and reduction trees. *Not DRAM*: traffic here does not contribute to




Figure 1: The GPU memory hierarchy. Cost per byte rises by roughly an order of magnitude per level descended; the data-movement problem is that most bytes live near the bottom while compute happens at the top. All DRAM traffic studied in this thesis is at the DRAM level and below.

the memory wall.

Local A per-thread region physically *in DRAM* but logically compiler scratch: when a kernel requires more temporaries than its register budget, the compiler *spills* them here (LDL/STL). Local addresses are interleaved across threads, so spill traffic is *perfectly coalesced by construction* a fact that resulted in this project’s most instructive measurement artifact (§4.3.3).

Constant, texture Small read-only parameter storage and image reads through dedicated texture units. Texture fetches do not appear in the address-trace instrument used here; the implications are discussed in §4.4.6.

2.1.3 The Memory Wall

Since the mid-1990s, processor arithmetic throughput has increased much more rapidly than improvements in memory bandwidth or latency the *memory wall* [3]. On modern hardware, a double-precision multiply costs picojoules, whereas fetching its operands from DRAM costs hundreds of picojoules to nanojoules. Crossing PCIe to host memory or storage is even more expensive. For low-arithmetic-intensity codes image pipelines being the canonical example the processor spends

most of its time waiting. The roofline model [11] formalizes this (§3.3.4): performance is bounded by $\min(\text{peak compute}, \text{AI} \times \text{peak bandwidth})$, where arithmetic intensity (AI) is FLOPs per DRAM byte. A kernel that falls under the sloped region of the roof cannot be improved by adding arithmetic units only by moving fewer bytes or reducing the distance they travel. The latter is the premise of memory-centric architecture.

2.1.4 Visual SLAM and cuVSLAM

Visual SLAM estimates a camera’s trajectory while building a map of the environment [12], [13]. The canonical pipeline consists of four stages:

1. **Front-end.** Each incoming frame is converted to grayscale; an *image pyramid* (the frame at successively halved resolutions) and gradient images are built. Distinctive corners are detected (Good Features To Track, using the Shi–Tomasi criterion) and tracked frame-to-frame with Lucas–Kanade optical flow. This stage runs on every frame at camera rate.
2. **Visual odometry.** From tracked features, the camera’s motion is estimated by solving a nonlinear least-squares problem.
3. **Bundle adjustment (BA).** Poses and 3D landmark positions over a window are refined jointly. Numerically, each iteration builds and solves a linear system (a Hessian reduced by the Schur complement to eliminate landmark variables); this thesis refers to that structure as `ba_linear_system` and measures precisely which kernels stream it.
4. **SLAM layer.** Selected frames become *keyframes*; their feature descriptors enter a *keyframe database*. When the camera revisits a mapped location, a database scan detects the *loop closure* and corrects the accumulated drift of the entire trajectory. Loop closure is rare (triggered on revisits) but is essential for consistent, long-session maps and its database is the only data structure that *grows* with session length, an architectural detail of central importance (§4.4.3).

The production implementation of the library is cuVSLAM [2]. It contains solvers

and a frontend running on the GPU, as well as a host-managed database of landmarks and keyframes based on LMDB (an embedded memory-mapped database). It supports multiple configurations including monocular, stereo, RGB-D (color and depth), and IMU-based inertial sensors. All of these may operate either in SLAM or odometry mode, asynchronously or synchronously. In this research, cuVSLAM is evaluated only as a test subject: the algorithmic implementations remain unchanged, the profiling instrumentation is optional and validated to have no impact on the output in Chapter 4.1.

The accuracy of a SLAM system is assessed against *ground truth* (GPS/INS, motion capture, or other tracker) according to two widely-used metrics [14], [15]. The first one is *absolute pose/trajectory error* (APE/ATE): the root-mean-square deviation of the position after optimal alignment (according to SE(3) for systems with known metric scale or Sim(3) for uncalibrated monocular systems [16]). The second one is *relative pose/trajectory error* (RPE/RTE): the cumulative drift over certain segments of the trajectory, measured as a percentage. Thanks to the relative error metric not requiring a ground truth in terms of absolute scale, it enables comparison between the large-scale automotive KITTI odometry dataset and small-scale room paths of TUM or EuRoC MAV.

2.1.5 Candidate Substrates for Memory-Centric Execution

This thesis evaluates each kernel’s suitability for five possible execution environments:

- **GPU (status quo):** Best when there is significant arithmetic per byte or when cache reuse is possible.
- **CPU/host:** Suitable for tiny, launch-overhead-dominated kernels with negligible GPU occupancy.
- **DRAM-PiM (near-bank):** Computation placed at DRAM banks leverages internal bandwidth several times greater than the external interface and eliminates the interconnect energy cost per byte. Commercial realizations include UPMEM’s DPU’s [8], Samsung’s HBM-PIM [9], and SK hynix’s GDDR6-AiM [10]. The ideal candidates are *streaming*, *bandwidth-bound*

tasks: large data, little arithmetic, no reuse.

- **Scatter-capable PiM:** Irregular gather/scatter operations defeat both caches and coalescing; near-memory engines capable of fine-grained random access (such as Tesseract and PIM-enabled instructions [17], [18]) represent the architectural solution when access patterns cannot be regularized in software.
- **Near-sensor SRAM / ISP:** Per-frame pixel processing can occur adjacent to the image sensor, consuming the stream before it ever reaches DRAM. The measured 1.68 MB-per-frame host-to-device upload (§4.2.7) is precisely the traffic this placement would eliminate.
- **In/near-storage processing:** For growing databases that are scanned episodically, the scan can be executed inside the storage device [6], [7], returning matches rather than streaming the entire store through host memory, PCIe, and DRAM.

2.1.6 The Measurement Instruments

No single tool provides a complete picture; the methodology combines four perspectives:

Nsight Systems (nsys): Records the timeline every kernel launch, its duration, and every PCIe copy with near-zero overhead. It provides the screening signal (identifying which kernels matter) and, via NVTX ranges, the kernel-to-pipeline-stage mapping.

Nsight Compute (ncu): Reads hardware performance counters per kernel by *replaying* each kernel multiple times. It is precise about *how much* (bytes, hit rates, stall cycles) but blind to *which addresses* are accessed. Replay is computationally expensive, so measurements are limited to selected windows and a curated metric set of roughly 15–30 counters; requesting the full set is impractical on small-memory GPUs.

NVBit [19]: Rewrites GPU machine code at load time to invoke a callback on every memory instruction, producing the exact per-warp address stream ground truth beneath ncu’s aggregates at 100–1000× slowdown. The tool used here (`mem_trace`,

extended with launch-window and kernel-name gating, plus an allocation-event sidecar) makes gigabyte-scale traces feasible and, when joined with the allocation journal of Chapter 4.4, enables attribution.

NVTX + TaggedAllocator: Assigns names to things. cuVSLAM includes NVTX stage annotations that are compiled out by default at two independent gates; this project’s instrumentation patch enables them, making the kernel-to-stage table a *measured* artifact rather than a guess based on kernel names. The TaggedAllocator (see Chapter 4.4) journals every GPU allocation along with its call stack, giving each buffer a durable data-structure name.

A standing epistemic ladder organizes these instruments: *counters reveal how much; traces show which addresses; the allocation journal identifies which data structure; NVTX indicates which pipeline stage*. Every claim in this thesis cites the highest available rung, and the two occasions where rungs disagreed became findings (§4.3.3, §4.4.1).

2.2 Related Work

2.2.1 Data-Movement Characterization Methodologies

The direct precursor to this thesis is DAMOV [1], which systematized the question: “Which functions of a workload are data-movement-bound, and would near-data processing help?” for CPUs. DAMOV’s three-step process screen functions that matter; characterize them with hardware counters plus architecture-independent locality analysis; and classify them into bottleneck families *derived from the data by clustering* is adopted here in full. Its central metric, the last-to-first miss ratio (LFMR: of the requests that miss the first cache, the fraction that also miss the last and reach DRAM), is retained in a redefined form. DAMOV’s validation discipline is also followed: thresholds are frozen and stress-tested on held-out inputs, an independent clustering algorithm reproduces the classes, and cache-size sweeps confirm class behavior. Chapter 3 enumerates, component by component, the GPU analogue of each DAMOV element and highlights the three places where GPUs required a redesign rather than a direct translation. Where DAMOV’s third step uses a simulator-based intervention (host vs. host-with-prefetcher vs. NDP

across a core-count sweep in ZSim/Ramulator), this thesis substitutes two *measured* interventions on silicon independent core and memory-clock scaling, and the reuse-distance cache-capacity curve and postpones the simulated axis to generated configuration (§3.7).

On the GPU side, the counter-metric curation employed here follows the NCU-based roofline methodology of Cao et al. [20], who demonstrated a disciplined path from Nsight Compute counters to roofline placement for GPU database kernels. This thesis reuses their metric taxonomy (extended with stall-reason and occupancy axes) and their practice of measuring, rather than quoting, roofline ceilings. The roofline model itself is due to Williams et al. [11]. Reuse-distance (stack-distance) analysis originates with Mattson et al. [21]; its use here as a *measured* hit-rate-versus-capacity curve from binary-instrumentation traces (§4.3.1) is the GPU translation of DAMOV’s architecture-independent locality step. One novel addition, not present in the CPU literature, is filtering by memory space, without which register-spill and scratchpad traffic can corrupt the locality signal (§4.3.3).

2.2.2 Processing-in-Memory and Near-Data Systems

The PiM literature spans over three decades; the modern wave is surveyed by Mutlu et al. [4] and Ghose et al. [5]. Architectures directly relevant to this thesis include Tesseract [17] (near-memory graph processing, which serves as the lineage for scatter-capable PiM), PIM-enabled instructions [18] (fine-grained offload of cache-defeating operations), TOM [22] (transparent offload of GPU code blocks to 3D-stacked memory, selected by compiler bandwidth-benefit analysis a methodologically close GPU-offload ancestor), and the workload-driven placement study of Boroumand et al. [23] for consumer workloads. Commercial near-bank silicon UPMEM DPUs [8], Samsung HBM-PIM [9], and SK hynix GDDR6-AiM [10] demonstrates that the substrates assigned in this thesis are realizable, not hypothetical. In-storage processing is grounded in Smart-SSD query offload [6] and Biscuit [7]. However, none of these works provide the *workload side* for Physical AI: a measured statement of which SLAM computations and data structures are best suited to each substrate. This is the gap addressed by this thesis.

2.2.3 SLAM Systems and Their Evaluation

The SLAM algorithmic literature is well established; Cadena et al. [12] provide a comprehensive survey. The systems used for accuracy validation in this thesis are cuVSLAM itself, as documented in NVIDIA’s technical report [2], and ORB-SLAM2/3, which serve as standard open-source references [24], [25]. Benchmarking SLAM *as a computational workload* rather than solely for accuracy was pioneered by SLAMBench [26], which profiled dense SLAM across CPU and GPU platforms at whole-pipeline granularity. This thesis differs in three key ways: it measures a *production* sparse V-SLAM stack rather than a research kernel; it achieves per-kernel and per-data-structure granularity; and its output is substrate placement, not a platform comparison. Trajectory-accuracy methodology follows the TUM benchmark conventions [14] and the evaluation tutorial by Zhang and Scaramuzza [15], incorporating Umeyama alignment [16]. The datasets used KITTI [27], EuRoC [28], TUM RGB-D [14], TUM-VI [29], and ICL-NUIM [30] are the standard ground-truth corpora in the community.

2.2.4 GPU Instrumentation

Binary instrumentation of GPU machine code is provided by NVBit [19], upon which this thesis’s address tracing is built (with two extensions: launch-window/kernel-name gating to limit trace volume, and an allocation-event sidecar that timestamps allocator activity using the trace’s own launch-id clock). GPU simulation and power modeling Accel-Sim [31] and AccelWattch [32] appear in this thesis only as the targets for generated configurations in the follow-on study (§7.1); in accordance with project policy, no simulated numbers are reported in the results chapters.

2.2.5 Positioning

In summary, relative to DAMOV, this thesis is the GPU re-derivation plus the extension from bottleneck class to substrate verdict. Relative to GPU profiling practice, it adds machine-independent locality and named data-structure attribution on top of counters. Relative to the PiM systems literature, it provides the missing measured workload analysis for a deployed Physical-AI perception stack.

And relative to SLAM benchmarking, it is the first study at the granularity where architectural decisions are actually made the kernel and the allocation.




Figure 2: The GPU-DAMOV pipeline. Capture is uninstrumented for timing; instrumented traces feed Attribute only. Every run also records energy, host I/O, and trajectory accuracy, so no measurement from a malfunctioning run enters the evidence base.

3 Methodology

This chapter defines the method. Section 3.1 presents the pipeline; §3.2 describes the instruments and workloads; §3.3 details every metric with its formula; §3.4 introduces the classifier and taxonomy; §3.5 covers substrate-verdict and fault mapping; §3.6 explains the placement model; and §3.7 provides the component-by-component parity argument that GPU-DAMOV is complete at DAMOV’s granularity (RQ1).

3.1 The Six-Step Pipeline

The method is a fixed pipeline (Figure 2); each stage consumes the previous one’s output, and the final stage emits a per-kernel substrate verdict.

Capture runs the workload under three profilers at locked clocks: nsys for timeline, ncu for counters (bounded launch windows, curated metric set), and NVBit for address traces (windowed). Each run records whole-run energy (GPU power sampled and integrated), host-side I/O (storage bytes read, memory-mapped page-ins, peak host RSS), and always, trajectory accuracy against ground truth ensuring no measurement from a malfunctioning run enters the evidence.

Screen ranks kernels by GPU time share; kernels below the attention floor are screened out, not classified (they receive class G0).

Classify passes each surviving kernel’s metrics through the ordered decision tree (§3.4).

Attribute joins address traces with the allocation journal to name each kernel’s traffic (see Chapter 4.4).

Validate stress-tests everything: threshold sensitivity, unsupervised clustering, cross-device agreement, a known-truth calibration suite, and profiler-neutrality checks (Chapters 4.1–4.2).

Verdict maps class $\times$ data-structure persistence $\times$ features to a substrate, with a named on-GPU fault when the verdict is “stay.”

3.2 Experimental Setup

3.2.1 Hardware

Testing is done primarily using an NVIDIA RTX 2000 Ada (**sm_89**): 22 SMs, 16 GB GDDR6 memory, 25 MB L2 cache, and 100 KB L1/shared memory per SM. To ensure consistency of results across different devices, the second, different in architecture device, GeForce MX450 (**sm_75**, 2 GB, 25 W), is used exclusively for cross-device validation (§4.2.4). The host machine is equipped with CUDA 12.9.1 and driver 575.64.05. The profiling is done using Nsight Compute 2025.2.1.3 and Nsight Systems 2025.3.2. NVBit 1.8 [19] is used for address tracing. To ensure reproducibility, this particular combination of hardware and software is carefully described in Appendix 9.2.

3.2.2 Locked Clocks and Measured Ceilings

GPU frequencies are locked at 1620 MHz (graphics) and 7001 MHz (memory) for all measurements, ensuring comparability. The effect is significant: the median run-to-run coefficient of variation ($\text{CoV} = \text{std. dev.} \div \text{mean}$) of kernel time drops from **49.6%** unlocked (thermally limited laptop), and 9.3% warmed, to **0.14%** locked. At 0.14%, run-to-run noise is far below every reported effect size.

Roofline ceilings are *measured*, not quoted: at locked clocks, sustained DRAM bandwidth is 205.0 ± 0.1 GB/s and FP32 throughput is 5445 ± 3 GFLOP/s (see artifact). All roofline and DRAM speed-of-light percentages in this thesis use these measured ceilings.

3.2.3 Workloads and Datasets

The testing dataset consists of 27 different sequences coming from four known ground-truth datasets: kilometre-scale automotive stereo data from KITTI odometry (00–10) [27]; room- and hall-scale drone stereo data with IMU from EuRoC MAV (MH01–05, V101–103, V201–203) [28]; handheld RGB-D captures from TUM RGB-D fr3 [14], including the `long_office_household` sequence; and the TUM-VI dataset for visual-inertial data [29]. Besides these, the ICL-NUIM dataset [30] is included for the accuracy-validation matrix (Chapter 4.1). The sequences cover a wide spectrum of testing conditions: indoor and outdoor settings, tight rooms and streets, and paths with and without loop closures (KITTI 01 and 04 used as baseline for loop-free tests). The testing pipeline configurations include SLAM vs. odometry-only, asynchronous vs. synchronous mode, CPU vs. GPU backend, and different types of sensors (inertial, monocular, RGB-D, and stereo). Using deterministic mutations, exactly 141 accuracy configurations and 192 profiling variants were generated, covering the entire testing matrix bit-for-bit.

3.2.4 Capture Discipline

Three rules are enforced: (1) *Traces are never used for timing*: instrumented runs are 5–1000 $\times$ slower; all timings are from uninstrumented, locked-clock runs. (2) *Windows must be audited*: windowed captures may miss sparsely-launched kernels; a coverage audit diffs each sequence’s full launch map against the captured set, gap-filling until all kernels are present (§4.4.2). (3) *Extensive before intensive*: time, bytes, and instruction counts are summed over all launches before computing ratios; averaging per-launch ratios would overweight short launches.

3.3 Metrics

Each metric reduces hardware counters or trace records to a single interpretable feature per kernel.

3.3.1 Speed-of-Light Utilizations

Nsight Compute reports each subsystem’s utilization as a percentage of its theoretical peak: *memory SoL*, *compute SoL*, and *DRAM SoL*. DRAM SoL $\geq 50\%$ is bandwidth saturation; compute SoL much above memory SoL is compute-boundedness.

3.3.2 LFMR (GPU Form)

$$\text{LFMR} = 1 - (\text{L2 hit rate}). \quad (1)$$

On a two-level cache, LFMR $\approx$ fraction of L2 accesses that go to DRAM. LFMR ≈ 1 : L2 not helping (PiM-favourable). LFMR ≤ 0.35 : L2 absorbs reuse; near-data placement forfeits a working cache.

3.3.3 Memory Intensity (MPKI)

$$\text{MPKI} = \frac{\text{DRAM bytes}/32}{\text{warp instructions}/1000} \quad (2)$$

The number of DRAM 32-byte-sector fetches per thousand *warp* instructions executed. (Note that the denominator for warp-level measurement is $32\times$ smaller than for thread-level; mixing them results in the classic off-by-a-factor-of-32 error.)

3.3.4 Arithmetic Intensity and the Roofline

$$\text{AI} = \frac{\text{FLOPs}}{\text{DRAM bytes}}, \quad \text{FLOPs} = \text{adds} + \text{mults} + 2 \times \text{FMAs}. \quad (3)$$

Attainable performance is $\min(P_{\text{peak}}, \text{AI} \times B_{\text{peak}})$ using the *measured* ceilings of §3.2.2 [11]. A kernel under the sloped segment is memory-bound; under the flat, compute-bound (Figure 3). The FLOP numerator selects the dominant operand (FP32 for cuVSLAM; FP16/INT for others), keeping the roofline meaningful across workloads.




Figure 3: Measured roofline (locked-clock ceilings) for representative TUM-office kernels. The dominant tracking kernels (`photometric`, `point_to_point`) sit deep under the sloped roof at low intensity (memory-wall territory) while the front-end convolutions and the BA reduction achieve a large fraction of the compute peak.

3.3.5 Occupancy and the Stall Taxonomy

Occupancy is the achieved fraction of maximum resident warps; below 25%, latency-hiding is starved. When a warp cannot issue, the hardware records why; the classifier uses the dominant reason: `long_scoreboard` (waiting for DRAM/L2), `short_scoreboard` (shared-memory), `lg/mio/tex throttle` (memory-pipe queues full), `wait` (dependency chains not memory), `barrier` (sync). The pair (stall, occupancy) distinguishes “memory-latency-bound” (G4) from “dependency-bound” (G7).

3.3.6 Coalescing

Sectors per request, from counters, with a trace-side check of sectors-per-warp-access from NVBit (§4.3.1); ≥ 8 marks a scattered gather.

3.3.7 Locality (Trace-Side)

See §4.3.1: reuse-distance CDFs, footprints, and inter-launch overlap, computed from global-space addresses only.

Table 2: Classifier thresholds (values as frozen in `analysis/classify.py`).

threshold	value	meaning
<code>sol_hi</code>	40 %	a speed-of-light value counted as “high”
<code>sol_ratio</code>	1.5×	margin by which compute SoL must exceed memory SoL for compute-boundedness
<code>dram_sat</code>	50 %	DRAM SoL treated as bandwidth saturation
<code>lfmr_hi</code> / <code>lfmr_lo</code>	0.40 / 0.35	L2 “not helping” / “earning its keep”
<code>sect_scatter</code>	8	sectors/request marking a scattered gather
<code>occ_low</code> / <code>occ_low_dep</code>	25 % / 30 %	occupancy below which latency (G4) / dependency (G7) stalls cannot be hidden

3.4 The Classifier and Taxonomy

3.4.1 Thresholds

Classifier thresholds (Table 2) are defined once in source and imported into all documentation and the dashboard, preventing drift. They are calibrated on the characterization corpus, then *frozen*; robustness is established by $\pm 25\%$ sensitivity stress (§3.4.3) and a known-truth calibration suite (§4.2.5).

3.4.2 The Ordered Rules

The first matching rule assigns the class:

1. **G5 compute-bound:** compute SoL $\geq 40\%$ and $\geq 1.5\times$ memory SoL.
2. **G6 on-chip-bound:** a shared-memory/MIO stall dominates while DRAM is not saturated.
3. **G1 bandwidth-bound:** DRAM SoL $\geq 50\%$.
4. **G2 coalescing-bound:** memory-limited and sectors/request ≥ 8 .
5. **G3 L2-reuse-bound:** memory-limited and LFMR < 0.35 .
6. **G4 latency-bound:** a memory stall dominates at occupancy $< 25\%$ with




Figure 4: The classifier as an ordered tree; the first satisfied test fixes the class. Orange leaves (G1, G2, G4) carry near-data affinity; grey leaves (G3, G5, G6, G7) stay on the GPU. Thresholds are those of Table 2, frozen in source.

DRAM unsaturated.

7. **G7 dependency-bound**: a dependency stall (`wait`, `barrier`, ...) dominates at low occupancy with neither memory nor compute high.
8. **G0 no-signal**: nothing dominant (tiny or launch-tax kernels; screened rather than architecturally interpreted).

Class G7 was not part of the original hypothesis; it emerged from cuVSLAM kernels that stall on instruction dependencies at low occupancy with memory idle. The taxonomy grew a class in response to the data; by design, DAMOV’s own classes were likewise derived. Figure 4 renders the same ordered rules as a tree: each internal node is one test, and the first satisfied branch fixes the class.

Table 3: The GPU bottleneck taxonomy and its default near-data affinity.

class	meaning	PiM affinity
G0	no dominant signal / sub-floor	n/a (candidate for CPU)
G1	DRAM bandwidth saturated	strong
G2	scattered access wastes bandwidth	conditional (scatter PiM or layout fix)
G3	L2 reuse effective	weak (keep the cache)
G4	memory-latency stalls, low occupancy	strong if cache-defeating
G5	compute-bound	none (keep on GPU)
G6	on-chip/shared-memory-bound	none
G7	dependency stalls, memory idle	none (fix occupancy first)

3.4.3 Sensitivity Analysis

Every kernel is re-classified with all thresholds scaled by $\pm 25\%$. Any kernel whose class changes under this perturbation is flagged as *borderline* and does not carry high confidence into the verdict stage; every headline kernel in this thesis is threshold-stable. Observed flips concentrate where physics predicts: at the L2-capacity crossover, where a working set near 25 MB moves between G1 and G3.

3.5 From Class to Verdict: Substrate and Fault Mapping

Class alone does not determine placement; it is joined with data-structure *persistence* (see Chapter 4.4). Persistence takes three values, measured from the allocation journal and stage map: *streaming* (per-frame data, consumed once), *hot-persistent* (resident state reused every frame, e.g., the BA linear system), and *cold-persistent* (session-lifetime state touched episodically, e.g., the keyframe database). Table 4 gives the mapping.

When the verdict is “stay on GPU,” the same evidence names the fault: spill-dominated traffic $\Rightarrow$ register-pressure/compiler fault (the traffic is scratch, not data); high sectors/request $\Rightarrow$ data-layout fault; dependency stalls at low occupancy $\Rightarrow$ launch-configuration fault; a flat reuse CDF $\Rightarrow$ structurally cache-immune (no cache-size remedy exists). This dual output substrate candidacy *or* a named fix makes the characterization actionable in both directions.

Table 4: Substrate verdict mapping.

condition	affinity	substrate
G1 $\wedge$ streaming	strong	near-sensor SRAM (consume before DRAM)
G1 $\wedge$ otherwise	strong	DRAM-PiM (near-bank)
G2	conditional	scatter-capable PiM, or data-layout fix first
G4 $\wedge$ cache-defeating	strong	near-memory compute
cold-persistent $\wedge$ large	strong	in/near-storage (ISP)
G3	weak	GPU (a larger cache wins)
G5 / G6 / G7	none	GPU
tiny (occ < 8%, < 1 ms)	none	CPU/host

3.6 The Placement Model

To bound the payoff of offloading without a simulator, a first-order analytical model is used: for a kernel with GPU time t and memory-bound fraction m , executed on a near-data substrate with internal-bandwidth advantage k and relative compute throughput c ,

$$t_{\text{pim}} = t \frac{1-m}{c} + t \frac{m}{k}. \quad (4)$$

A kernel is offloaded only if its affinity is admitted by the scenario and $t_{\text{pim}} < t$. Two scenarios bound the claim: *conservative* ($k=4$, $c=0.5$, strong-affinity only) and *moderate* ($k=8$, $c=0.75$, strong + conditional). An energy variant applies the same split to the measured energy baseline. A standing rule constrains all of this: model-derived and (future) simulator-derived numbers are reported only as *deltas against the measured GPU baseline*, never as absolute performance claims.

3.7 DAMOV Parity: The Completeness Argument

RQ1 asks whether every DAMOV component has a GPU analogue here. The correspondence, element by element, is as follows:

- *Step 1, screening.* **DAMOV:** VTune top-down, retaining functions with >30% memory-bound share and $\geq 3\%$ of cycles. **Here:** nsys time-share plus the ncu SoL pair; sub-floor kernels are screened (G0), not classified.
- *Step 2, locality.* **DAMOV:** Architecture-independent spatial and tempo-

ral locality from CPU address streams, clustered. **Here:** NVBit per-warp traces $\rightarrow$ reuse-distance CDFs (temporal) and sectors-per-warp coalescing (the GPU’s spatial analogue), with *memory-space filtering* to prevent compiler scratch from polluting locality (§4.3.3).

- *Step 3, intervention.* **DAMOV:** Simulated host / host+prefetch / NDP across a 1–256-core sweep, with classes defined by the response. **Here:** Two measured axes replace the unavailable core sweep (a GPU grid is compiled into the launch; SM count is fixed): (a) *Clock-domain sweep* core and memory clocks are scaled independently (1620 $\rightarrow$ 810 MHz core = $0.5\times$ compute; 7001 $\rightarrow$ 5001 = $0.71\times$ DRAM) yields falsifiable per-class predictions (G1/G2 track the memory clock, G3/G5/G6/G7 the core clock, G4 neither strongly); (b) *Measured* reuse-distance hit-CDF directly answers the cache-capacity question for each kernel, from 64 KiB to 48 MiB, where DAMOV used aggregate sweeps. The third, simulated axis (NDP configurations for Accel-Sim) is generated from the verdicts and deferred to the follow-on study.
- *Classification.* **DAMOV:** Six classes from a threshold tree over temporal locality, LFMR trend, MPKI, and AI. **Here:** Eight classes (G0–G7) from the ordered tree of §3.4; the deviations (a coalescing class, a G4/G7 occupancy split, and no L3-contention class because L2 is the last level) are each necessitated by GPU-specific differences.
- *Robustness.* **DAMOV:** Frozen thresholds validated on 100 held-out functions (97% correct); cache-size sweeps; an independent clustering algorithm. **Here:** $\pm 25\%$ sensitivity stress; a known-truth *calibration suite* of eight archetype kernels, classified blind (§4.2.5); cross-device agreement across two microarchitectures (§4.2.4); and *two* independent clustering algorithms (k-means and Ward hierarchical) recovering the taxonomy (§4.2.3).

Where DAMOV’s breadth was 144 applications, this study is deliberately deep rather than wide one production application, 49 kernels, 27 sequences, 192 configuration variants with architectural breadth provided by the harness’s adapter interface, capable of running arbitrary GPU workloads through the identical pipeline. The calibration archetypes serve as the suite-artifact analogue.

4 Results

4.1 Measurement Validity: Accuracy, Neutrality, and Noise

A characterization inherits the validity of its measurements. This chapter establishes, in order: that the system under measurement was computing *correct* trajectories (§4.1.1); that observing it did not change its behaviour (§4.1.2); and that residual run-to-run variation is understood and bounded (§4.1.3). Together, these answer RQ5 and constitute finding **F13** in the evidence chain.

4.1.1 Accuracy Validation: The 141-Configuration Matrix

Every combination of dataset, sensor modality, and pipeline mode that the on-disk corpora support was instantiated as a configuration 141 in total. Each was run to completion and scored against ground truth using the vendor’s own metrics (RMSE APE after alignment; segment-relative RTE; rotational RE). A convergence gate (segment-relative translational drift below 5%, a scale-independent criterion) separates runs that track from those that diverge; every excluded run is itemized with a stated cause rather than silently dropped.

Table 5 compares the converged mode averages against the cuVSLAM technical report [2]. Two findings support this chapter. First, *the modes used for the entire memory characterization stereo and RGB-D SLAM reproduce the published accuracy*: EuRoC stereo matches to within millimetres, TUM RGB-D is better than published, and KITTI matches on the definition-stable per-segment drift metric (0.82% at the 500 m segment versus 0.85% in the leaderboard; absolute APE over kilometre paths is scale-sensitive and reported as directional only). Second, every discrepancy in other modes has a named, verified cause, none of which affects the characterized configurations.

4.1.1.1 Discrepancies and Their Causes Inertial mode is under-tuned by configuration, not by cuVSLAM. Only a minority of EuRoC inertial runs converge, and diagnostically enabling the IMU makes accuracy *worse* than pure stereo on the same sequences. This inversion is the classic signature of mis-scaled IMU noise densities or camera–IMU extrinsics: the estimator over-trusts a mis-

Table 5: Converged mode-average accuracy vs. the cuVSLAM technical report (RMSE APE in metres, odometry/SLAM). The characterization uses the stereo and RGB-D modes.

mode	ours	published	verdict
EuRoC stereo	0.114 / 0.051	0.13 / 0.054	millimetre-level match
TUM RGB-D	0.060 / 0.047	0.11 / 0.065	better than published
ICL-NUIM (mono-depth)	0.099 / 0.136	0.026	same order (cm)
KITTI stereo	4.01 / 3.63	3.00 / 1.98	segment drift matches (0.82% vs 0.85%)
EuRoC stereo-inertial	0.401 / 0.328	0.19 / 0.13	under-tuned IMU (disclosed)

modelled sensor. The vendor calibrates these per dataset; the matrix used a generic inertial configuration. This defect is confined to a mode never used in memory characterization, and the remedy (populating per-sequence IMU parameters from the datasets’ calibration files) is configuration work, tracked in the roadmap.

The hardest EuRoC sequence diverges in every mode. V2_03 combines aggressive motion with heavy motion blur; the vendor’s report publishes it only under stereo-inertial and excludes it from stereo averages. The same exclusion is applied here, with the sequence retained in per-sequence tables.

TUM-VI raw fisheye is invalid input, not a tracking failure. TUM-VI’s $\sim 195^\circ$ fisheye exceeds the pinhole-undistortion field of view supported by the library ($< 180^\circ$); using raw fisheye leads to deterministic, kilometre-scale divergence. The required preprocessing (undistort to pinhole with edge masks) is documented in the dataset notes; runs without it are classified as invalid input and excluded.

Monocular needs scale-free alignment. Monocular SLAM recovers shape but not metric scale; scoring it with $SE(3)$ alignment conflates scale error with drift. Monocular rows are therefore excluded from $SE(3)$ comparisons pending $Sim(3)$ evaluation [16] an evaluator-metric choice, not a tracking result.

4.1.2 Profiling is Accuracy-Neutral

The deepest reflexivity threat is that instrumentation changes the system it observes. Three experiments rule this out:

Instrumented build equals baseline build (QoR). The TaggedAllocator instrumentation (Chapter 4.4) lives in a rebuilt cuVSLAM binary. Re-running representative configurations on both builds shows the instrumented build’s trajectories are equal to the baseline’s bit-identical on EuRoC, and within run-to-run nondeterminism elsewhere (largest observed difference 0.24 m over a kilometre-scale KITTI path, within the sequence’s own re-run scatter; see §4.1.3). Journaling is gated by an environment variable and affects only allocation paths, which execute at initialization, not per frame.

Each profiler is trajectory-neutral. Running the same configuration plain and under each instrument shows: nsys and ncu produce *bit-identical* trajectories on deterministic (synchronous) modes, and NVBit-traced runs agree within ~ 2 mm APE the instrumentation slows execution greatly but does not alter the computation.

Neutrality holds across the full feature matrix. A 192-variant campaign (every accuracy configuration profiled as-is, plus finer feature toggles sync/async/CPU SLAM, motion-model and denoising toggles, landmark export, multi-camera precision, unrectified input, depth-stereo tracking on representative sequences per modality) compares profiled vs. plain APE with tolerance max(5 cm, 5%): **166/192 pass outright**, with 121 bit-identical. All 26 flagged variants were investigated and traced to benign causes mainly nondeterministic scatter from long loop-closing sequences (§4.1.3) and toggles that legitimately change computation (e.g., CPU-SLAM changes the back-end and hence the trajectory, independent of profiling); a re-check against plain re-runs disproved the only suspected genuine interference case.

4.1.3 Nondeterminism, Characterized

GPU SLAM is not perfectly repeatable: floating-point reductions reorder across runs, and asynchronous modes add scheduling freedom. Rather than averaging this away, the campaign quantified it. On short and mid-length sequences, repeated runs are bit-identical in synchronous modes. On long loop-closing sequences (kitti00 as canonical), the final trajectory is *bimodal*: loop-closure acceptance is

threshold-sensitive, and a borderline loop either fires or not, moving whole-path APE by tenths of a metre. The critical control: *plain, uninstrumented re-runs scatter by the same amount and between the same modes*. This scatter is a property of the workload, not of the profilers an important result for hardware studies (single-run accuracy comparisons on such sequences are meaningless at sub-scatter resolution).

4.1.4 Summary of the Trust Chain

The system under measurement reproduces vendor-published accuracy in characterized modes (141-run matrix); the instrumentation is output-neutral (bit-identical where the workload is deterministic, and within measured workload scatter elsewhere, across 192 variants); the noise floor of all timing is CoV 0.14% at locked clocks (§3.2.2); and every classification threshold is sensitivity-stressed (§3.4.3). Every number in Chapters 4.2–4.5 rests on this chain.

4.2 Kernel-Level Characterization

This chapter presents the classification of cuVSLAM’s 49 GPU kernels across the 27-sequence corpus and subjects that classification to four independent validations: input stability (§4.2.2), unsupervised recovery by two clustering algorithms (§4.2.3), cross-microarchitecture agreement (§4.2.4), and a known-truth calibration suite (§4.2.5). It closes with measurement of host–device transfers at the system edge (§4.2.7).

4.2.1 What the Screen Reveals

The nsys timeline over full sequences shows the expected structure of a frame-synchronous perception pipeline: a small set of front-end kernels launched every frame (image cast, Gaussian pyramid scaling, gradient convolutions, Lucas–Kanade tracking, and the matcher pair) accounts for the bulk of launches (~ 80 per frame; $\sim 200,000$ launches over a 2,500-frame sequence). The bundle-adjustment solver family (`sba::*`, plus cuSOLVER helper kernels) forms a second tier, while SLAM-layer keyframe kernels (`st_build_cache`, `st_track_with_cache`) launch

sparsely per keyframe and per loop-closure scan, respectively. The NVTX stage table (measured, not name-matched; §2.1.6) confirms stage membership: `st_track_with_cache` executes inside the loop-closure-and-optimization range, `st_build_cache` within keyframe ingestion the measured anchor for every loop-closure claim that follows.

4.2.2 Class Stability Across Inputs

The classification model presented in §3.4 was run independently on each sequence, and then aggregated over the entire dataset. The major conclusion is that there is **91% of agreement in the modal class**: out of 49 kernels, 24 belong to the exact same bottleneck class for all 27 sequences, while 42 belong to the same class in at least 80% of them. The bottleneck class is, to the first order, *the property of the kernel, not the input*.

The residual flips are not noise; they are physics. They cluster around kernels whose working set sits near the 25 MB L2 capacity: on small indoor sequences the set fits (L2-reuse behaviour, G3), on street-scale it spills (bandwidth behaviour, G1). The sensitivity analysis (§3.4.3) flags exactly these kernels as borderline, and the substrate synthesis (Chapter 4.5) treats their verdicts as workload-conditional. 25 of 49 verdicts flip with workload scale, always in the physically expected direction.

4.2.3 The Taxonomy Is Discovered, Not Asserted

The decision tree assigns labels; whether the *classes themselves* are real structure in the data is tested without labels. Figure 5 first shows how the 49 characterized kernels distribute across the classes. Pooling every kernel’s feature vector (memory/compute/DRAM speed-of-light, LFMR, occupancy, sectors-per-request, and the two dominant stall shares) across sequences and running k-means over a range of cluster counts, the silhouette criterion selects $k = 7\text{--}8$ matching the seven populated classes plus the screened remainder, with cluster purity 0.68 and adjusted Rand index 0.30 against the tree labels. Using a completely different algorithm Ward hierarchical clustering [33], cut at $k=8$ the same agreement is reproduced: **ARI 0.31, purity 0.675**. Two unrelated algorithms, given no thresholds and




Figure 5: Distribution of the 49 characterized kernels across the eight classes (campaign modal class). The near-data-affine classes (G1, G2, G4; orange leaves in Figure 4) account for 24 kernels. G6 is unpopulated in cuVSLAM (no kernel is on-chip-bound with DRAM idle), and G7, the class that emerged from the data, is the single largest, capturing the dependency-stalled tracking and solver kernels.

no labels, find the same ~ 8 -way structure encoded by the tree. The tree should be read as the *labelling* of natural structure and the clustering as its *independent validation* precisely DAMOV’s discipline, reproduced on GPU data.

4.2.4 Cross-Microarchitecture Agreement

If the classes captured artifacts of one GPU, they would not persist across microarchitectures. The same workload (TUM office) was characterized independently on two deliberately different devices: the `sm_89` workstation and an `sm_75` laptop with a quarter the memory and a third the power budget. **36 of 45 signal kernels (80%) receive the same class on both** (76.6% including the launch-tax G0 tier). The honest caveat: the two devices differ in memory system as well as core generation, so some disagreements are physically real L2-capacity crossovers rather than classification error making 80% a conservative floor for device-independence of the classification. This mirrors DAMOV’s core-type-independence check, under a stricter test.

4.2.5 The Calibration Suite: Classifying Known Truth

DAMOV validated frozen thresholds on held-out functions with known behaviour. Here, a *calibration suite* of eight archetype kernels is used, each with a known

ground-truth class by construction: a streaming triad (G1), a random gather (G2), an L2-resident sweep (G3), a coalesced pointer-chase (G4), an FMA polynomial (G5), a bank-conflicted shared-memory kernel (G6), a single-warp dependency chain (G7), and a sub-screen tiny kernel (G0). These are compiled and pushed through the identical capture-and-classify pipeline, blind. The first pass classified five of eight as designed; analysis of the three misses showed each to be a defect in the *archetype*, not the classifier: a per-lane-random pointer-chase is both scattered and latency-bound (G2+G4); a barrier-heavy shared-memory kernel presents as dependency-stalled rather than G6; and a sub-floor kernel is legitimately *screened*, as per the pipeline’s Step-1 semantics. Each archetype was corrected to embody a single class *with the classifier untouched throughout* and the suite is retained in the artifact as a permanent regression harness: any future change to thresholds or metrics must re-classify known truth correctly.

4.2.6 Near-Data Affinity of the Workload

Rolling the per-kernel classes up by measured GPU time: across the 27-sequence campaign, **21% of GPU time is in strong-affinity classes, 40% in conditional-affinity classes, and 28% in no-affinity classes** (the remainder are sub-floor). Thus, roughly **61% of GPU time carries near-data affinity** under the taxonomy’s default mapping, rising to $\sim 72\%$ time-weighted in deep single-sequence studies where the loop-closure path is exercised densely. This is the quantity the placement model in Chapter 4.5 prices. It is reported here with its composition, not as a headline slogan, since the conditional 40% is exactly the portion whose verdict depends on the scatter-PiM-vs-layout-fix question addressed in the locality and attribution chapters.

4.2.7 The System Edge: Host–Device Transfers

Independent of any kernel’s internals, the timeline measures the pipeline’s boundary traffic: explicit host-to-device and device-to-host copies account for **41% of total kernel time** on the RGB-D workload, and the per-frame sensor upload is **1.68 MB/frame** the camera image and depth crossing PCIe into GPU DRAM before any kernel touches them. This is the single most direct near-sensor argu-

ment in the thesis: it is traffic that exists only because the pixels are *born* on the wrong side of the bus, and no on-GPU optimization can remove it. The host-side companion measurements (66 MB of storage ingestion per run; 708 MB peak host memory) are detailed in the synthesis chapter, where the boundary picture is assembled end-to-end.

4.3 Machine-Independent Locality Analysis

Counters summarize; address traces prove. This chapter reports the DAMOV-style locality analysis computed from NVBit per-warp address traces: the metrics (§4.3.1), the front-end result (§4.3.2), the loop-closure-scan result and the self-correction that produced it (§4.3.3), and the session-scale growth measurement (§4.3.4).

4.3.1 Trace-Side Metrics

From each kernel’s global-space address stream, per launch, four metrics are computed:

- **Footprint:** Number of distinct 32-byte sectors touched, $\times 32$ B exact working-set size, measured directly rather than inferred.
- **Reuse-distance CDF:** For each access, the count of distinct sectors touched since that sector’s previous access (LRU stack distance [21]). The CDF, evaluated at capacities from 64 KiB to 48 MiB, predicts the hit rate of any cache of that size, independent of policy details. A flat curve means no cache size helps: misses are compulsory.
- **Coalescing, trace-exact:** Sectors touched per warp access (1–32), plus active lanes per access, which separates true scatter from divergence-induced apparent scatter.
- **Inter-launch overlap:** Jaccard similarity $|A \cap B| / |A \cup B|$ between consecutive launches’ footprints indicates whether a recurring kernel re-touches the same data (cacheable across launches) or migrates.




Figure 6: Predicted hit rate versus cache capacity, from the measured reuse-distance CDF. The front-end curve is essentially flat across a three-order-of-magnitude capacity sweep: no realizable cache changes its miss rate, because its misses are compulsory first-touches of streaming pixel data. A cache-friendly kernel (dashed) is shown for contrast.

One rule governs all of these: *memory-space filtering*. Only global-space instructions (LDG/STG) enter locality analysis; shared-memory (LDS/STS) and local-memory (LDL/STL, register spill) records are bucketed separately. This rule was not a precaution rather, it was forced by the measurement failure documented in §4.3.3, and is now the methodology’s most exportable lesson.

4.3.2 The Front-End Is Cache-Immune Streaming

For every per-frame front-end kernel (image cast, pyramid scaling, gradient convolutions, LK tracking), the reuse-distance CDF is **flat from 64 KiB to 48 MiB** (Figure 6): the hit rate a 64 KiB cache would achieve is the hit rate a 48 MiB cache would achieve, since all reuse is at tiny distances (caught by any cache) and everything else is a first touch of streaming pixel data (caught by none). Global-space accesses run at ~ 4.0 sectors per warp access, fully coalesced, with 32.0 active lanes.

The architectural reading is sharp: *no realizable cache helps this tier* its misses are compulsory. Bigger caches are the wrong axis entirely; the only remedies are moving fewer bytes (algorithmic) or consuming the stream before it reaches DRAM (near-sensor placement, §4.2.7). Early counter-only readings suggested one

convolution kernel was a scattered "data-layout target" (15.4 sectors per access) with high reuse; space filtering dissolved both signals these were shared-memory tile traffic, on-chip and irrelevant to DRAM. After correction, the front-end is *uniformly* stream-clean, and the correct optimization target is on-chip, not layout.

4.3.3 The Loop-Closure Scan: A Self-Correction on Record

The loop-closure scan kernel `st_track_with_cache` produced the project's most instructive measurement arc, reported in full for its methodological value.

Round one: counters vs. trace. Nsight Compute counters indicated the kernel was a scattered gather 18–30 sectors per request. The first NVBit trace analysis indicated the opposite: 2.1 sectors per warp access, 99.4% of accesses touching ≤ 4 sectors apparently tightly coalesced streaming. Taking the trace as ground truth, the project's interim report concluded the counter metric was a proxy artifact "the trace overturns the counters."

Round two: the join exposes the confound. When the attribution join (Chapter 4.4) matched trace addresses against live allocations, 88–98% of most kernels' trace records matched *no allocation at all* impossible for real data traffic. The explanation: the tracer records *every* memory space, and for `st_track_with_cache`, **92.6% of all recorded accesses were register-spill traffic** in local space which is *perfectly coalesced by construction* (§2.1.2) drowning the actual data accesses in the blended statistics.

Round three: space-filtered re-derivation. Filtering to global space only and re-deriving every table: `st_track_with_cache`'s *data* accesses are **23.4–30.0 sectors per warp access, with only 2.3–6.0% of accesses touching ≤ 4 sectors** a scattered gather, *exactly what the counters had said*. The interim conclusion was withdrawn in a dated, in-place correction; the classifier's G2 label stands; and the two instruments now agree. Table 6 juxtaposes the blended and corrected views.

Three durable lessons:

- *Methodological*: unfiltered GPU address traces are dominated by spill and scratchpad records; any locality or attribution analysis that does not filter

Table 6: The `st_track_with_cache` correction: blended-space analysis vs. space-filtered re-derivation (TUM room-scale / KITTI street-scale).

metric	blended (all spaces)	global-only	local-only (spill)
sectors / warp access	2.1	23.4 / 30.0	2.0
accesses ≤ 4 sectors	99.4 %	6.0 / 2.3 %	100 %
share of recorded accesses	100 %	$\sim 5\text{--}7\%$	$\sim 92.6\%$
per-scan footprint (MB)	0.467 / 1.096	0.464 / 1.093	~ 0.003
inter-launch Jaccard	0.67 / 0.90	0.67 / 0.90	1.0

by memory space is unsound a warning directly applicable to other groups using binary-instrumentation traces.

- *Evidential*: two independent instruments agreeing *after* a documented reconciliation is stronger support than either alone; the disagreement was the pipeline catching its own error.
- *Architectural*: the kernel’s DRAM volume and its DRAM *pattern* have different sources volume is dominated by compiler spill (a register-file/occupancy trade, fixable on-GPU), while the data-side pattern is a genuine scattered gather over the keyframe descriptors (a scatter-PiM or layout question). Only the attribution of Chapter 4.4 could have separated them.

4.3.4 Session-Scale Growth and Migration

Two properties of the scan survive the correction untouched (the spill window is only $\sim 3\text{KB}$ of the footprint). First, the per-scan working set *grows with the map*: 0.46 MB per scan at room scale (TUM) to 1.09 MB at street scale (KITTI). Second, it *migrates*: consecutive scans overlap with Jaccard 0.67 (room) to 0.90 (street) 10–33% of the touched set turns over between scans. Each individual scan is nearly L2-resident; the union over a session the entire keyframe database, scanned incrementally, growing without bound with distance travelled is what no cache holds. This is the measured locality half of the in-storage case; the allocation half (the database’s GPU-resident state is a fixed buffer while the growing store lives host-side) is established in Chapter 4.4.

4.4 Data-Structure Attribution: Naming Every Byte

A kernel-level claim (“this kernel is memory-bound”) is not yet an architect’s claim (“the keyframe database belongs near storage”). Bridging that gap requires knowing, for each byte of DRAM traffic, *which named data structure it belongs to*. This chapter presents the TaggedAllocator attribution: the instrument (§4.4.1), the coverage discipline (§4.4.2), the static-budget finding (§4.4.3), the per-kernel attribution results (§4.4.4), the register-spill finding (§4.4.5), and the bounded residuals (§4.4.6).

4.4.1 The Three-Layer Instrument

Attribution is a chain of fragile assumptions (address $\rightarrow$ allocation $\rightarrow$ owning code $\rightarrow$ data-structure name), so the design uses three independent layers for cross-checking:

Layer 1: source-side journal (semantic layer). Every GPU allocation in cuVSLAM flows through allocation wrappers; an injected journal records, for each allocation and free, the pointer, size, wrapper kind, and a host *call stack*. Resolved offline against debug symbols, the innermost non-plumbing frame names the owning data structure (e.g., the Schur-complement solver’s constructor $\Rightarrow$ `ba_linear_system`). A fixed tag vocabulary (`images_raw`, `pyramid_levels`, etc.) keeps names stable across runs and sequences. Pinned host buffers (device-visible under unified addressing) are journaled with a distinguishing suffix, as their traffic crosses PCIe rather than DRAM.

Layer 2: driver-side sidecar (temporal layer). The NVBit tool independently logs every driver-level allocation and free, stamped with the current kernel-launch counter, the same clock the address trace uses. This provides allocation lifetimes and catches allocations that bypass wrappers (e.g., math libraries’ scratch), which are tagged as driver-internal rather than silently mislabeled.

Layer 3: the join. A streaming pass over the address trace maintains the live-allocation set in launch order and resolves each global-space access to its containing allocation’s tag; shared-space and local-space records are bucketed as on-chip and spill, respectively (§4.3.3). Anything unresolved lands in an explicit `unmapped`

bucket, the join’s own error bar, reported and bounded.

4.4.2 Coverage: Auditing the Windows

Traces are captured in windows, and sparse kernels can drift out of a window between runs (launch indices shift with keyframe timing). The pipeline treats coverage as a proof obligation: after each campaign, an audit diffs every sequence’s *full* launch map (from a cheap uninstrumented pass) against the union of attributed kernels, and a planner captures gap-fill windows anchored on dense clusters of missing kernels’ launches (clusters are robust; isolated occurrences drift). Across the 27-sequence campaign this loop converged to **zero missing kernels**, with all audit artifacts committed. Without this discipline, three sequences would have silently lacked the loop-closure scan and one the keyframe-refresh cluster, precisely the architecturally interesting kernels.

4.4.3 The GPU Memory Budget Is Static

The journal’s first result is structural. Over a full 2,500-frame loop-closing session, cuVSLAM performs **240 device allocations totalling 108.65 MB all at initialization, none freed mid-run and the identical 240 in a 30-frame run**. The budget by tag (Table 7): images and pyramids ~ 71 MB; the BA linear system 24.3 MB (plus a 15.4 MB pinned-host mirror); and most decisively a **fixed 6.7 MB keyframe-descriptor buffer**. The keyframe *database* (the only data structure that grows with session length) does not grow on the GPU at all: its growing store lives host-side in the memory-mapped database, confirmed by direct host measurement (708 MB peak host memory against the static GPU budget; §4.5.5). The in-storage verdict of Chapter 4.5 rests on this measured division of labour, not on inference from kernel behaviour.

4.4.4 Attribution Results

The join resolves **274 of 274** journaled allocations (240 device + 34 pinned-host) with zero unknowns. Across the full 27-sequence campaign, **48 of 49 kernels receive a unanimous top data-structure tag in every sequence they appear in** the attribution analogue of the 91% class consistency of §4.2.2, and the

Table 7: The static GPU allocation budget (full session; peak = total, nothing freed mid-run).

tag	allocations	peak live (MB)
images_raw	92	39.97
pyramid_levels	99	31.15
ba_linear_system	31	24.26
keyframe_descriptors	2	6.73 (fixed)
feature_tracks	8	4.27
depth_pyramid	3	2.21
icp_state	4	0.07
device total	240	108.65
pinned host (device-visible)	34	17.14

license to make per-kernel data-structure claims input-independently. Selected rows (steady-state window; shares are sector-weighted):

- `st_track_with_cache` (loop-closure scan): **92.6% register spill**, with the global remainder on `keyframe_descriptors`; unmapped <1%.
- `st_build_cache` (keyframe ingest): 93% `keyframe_descriptors`.
- `sba::reduced_system_stage_2` (solver core): **96.9% ba_linear_system** one named, pre-sized, contiguous structure carrying essentially all of the solver’s DRAM traffic.
- Front-end compute kernels (convolutions, feature selection, sorting): 83–98% of accesses are shared-memory tiles (on-chip; no DRAM claim), with the global residue falling on the pyramid, image, and track structures the pipeline stage implies.

Mode differences surface exactly where expected: stereo configurations journal no depth or ICP structures; loop-free KITTI sequences show `st_track_with_cache` launches only when their trajectories genuinely close loops the attribution reproduces workload semantics, itself a correctness check.




Figure 7: Memory-space attribution for three representative kernels. The loop-closure scan `st_track_with_cache` spends 92.6% of its DRAM volume on register spill (a compiler artifact, not data) while its true data gather is the small remainder on `keyframe_descriptors`. By contrast the BA solver core and the keyframe-ingest kernel spend almost all of their traffic on a single named structure.

4.4.5 The Register-Spill Finding

The single most consequential attribution row is `st_track_with_cache`’s (Figure 7): the kernel motivating the near-data discussion spends **92.6% of its DRAM volume on register spill** compiler-generated scratch traffic, invisible to any counter-only methodology.

The finding cuts both ways. For *this* workload: the loop-closure scan’s bandwidth bill is mostly a register-pressure artifact (per-thread descriptor state exceeds register budget), so a register-file/occupancy trade or spill-aware compilation attacks the volume, while the residual the true database gather (§4.3.3) is the scatter-PiM candidate. For the *methodology*: spill is a first-class DRAM citizen that only per-space attribution can expose; a caveat applies (register allocation is a code-gen decision for this compiler/target; taxonomy transfers, spill *quantity* must be re-derived per target, and cross-target validation is scheduled work).

4.4.6 Bounded Residuals

The join’s honesty budget: for most kernels, the unmapped share is below 1–7%. Two systematic residuals remain, both bounded and explained. First, one solver-assembly kernel carries ~40% unmapped traffic *identically across all 27*

sequences the signature of static module memory (device-side globals allocated by the module loader, invisible to both journal layers), not a lurking data structure; analysis names the two kernels (`sba::build_full_system`, `lk_track`) that carry 83% of all unmapped traffic, making the residual a prioritized target for a planned kernel-argument correlation layer. Second, texture-path reads (the LK tracker samples pyramids through texture units) never appear in the address trace, so image-structure traffic is a stated lower bound, with counter-side texture totals corroborating the affected kernels. Neither residual moves any verdict: affected kernels’ classes rest on counters and stage evidence, independent of the unmapped share.

4.5 Synthesis: Substrate Verdicts and Measured Baselines

This chapter assembles the evidence from Chapters 4.2–4.4 into the thesis’s deliverable: a per-tier substrate placement with measurement behind each cell (§4.5.1), the dynamics of the verdicts across workloads (§4.5.2), the analytical placement bound (§4.5.3), and the measured energy and host-side baselines for follow-on architecture studies (§4.5.4, §4.5.5).

4.5.1 The Three-Way Split

cuVSLAM’s memory behaviour resolves into three tiers, each with a distinct substrate answer (Figure 8 and Table 8).

Each row is anchored by an NVTX stage, a bottleneck class, a locality curve, and a named allocation. The rows are *different* which is the finding. A single “SLAM accelerator” would target three mutually incompatible bottlenecks; the measured answer is heterogeneous placement.

4.5.2 Verdict Dynamics Across Workloads

Applying the verdict mapping (§3.5) per sequence: 24 of 49 kernels receive the same verdict unanimously; **25 of 49 flip with workload scale**, driven almost entirely by DRAM speed-of-light saturation as working sets grow from room to street scale the expected physical driver, not classifier noise. Time-weighted, **~72% of GPU**




Figure 8: *cuVSLAM* resolves into three tiers with distinct substrate verdicts. The streaming front-end wants near-sensor consumption; the compute-efficient bundle-adjustment solver stays on the GPU; the scattered, session-growing loop-closure back-end is the near-data (PiM-scatter and in-storage) candidate. The verdicts are per-data-structure, not per-application.

time is offload-eligible in deep studies (61% across the breadth campaign; §4.2.6). The practical consequence: placement for this workload class is *scale-conditional*, and a deployment at street scale crosses into near-data-favourable territory that a desk-scale evaluation would miss.

4.5.3 The Analytical Placement Bound

Pricing the eligible time with the placement model of §3.6: under the *conservative* scenario ($k=4$, $c=0.5$, strong-affinity only), **16 of 49 kernels** exceed the offload threshold; under the *moderate* scenario ($k=8$, $c=0.75$, strong + conditional), **26 of 49**. These are candidacy bounds not performance claims. The model’s k and c are scenario parameters; per the standing rule, follow-on simulation reports deltas against the measured baseline, not absolutes. What the bound shows: the opportunity survives pessimistic assumptions the offload set is non-empty even at $k=4$ with half-rate compute.

Table 8: The three-way substrate split, with measurement backing each claim.

tier	measured evidence	substrate consequence
Streaming front-end (images, pyramids, gradients, tracking)	reuse-CDF flat 64 KiB→48 MiB (compulsory misses); fully coalesced at ~ 4 sectors/access; 1.68 MB/frame sensor upload; 41% of kernel time in PCIe copies	near-sensor SRAM / ISP: consume pixels before DRAM; caches cannot help; upload disappears by construction
Hot-persistent solver (BA linear system)	96.9% of solver-core traffic on one pre-sized 24.3 MB structure (+15.4 MB pinned mirror); compute-leaning classes; L2-effective where resident	keep on GPU at this scale (bank-level PiM streaming is a candidate if problem scale outgrows L2)
Cold-persistent back-end (keyframe database)	GPU-resident state fixed at 6.7 MB while the store grows host-side (708 MB peak host RSS); scan is a scattered gather (23–30 sectors/access) whose footprint grows 0.46→1.09 MB room→street with 10–33% turnover per scan; 92.6% of scan DRAM volume is register spill	in/near-storage for the growing store; scatter-capable PiM or layout fix for the resident gather; register-file/spill SRAM for the volume

4.5.4 The Measured Energy Baseline

Whole-run GPU energy, measured by integrating sampled board power at locked clocks: **34.67 J for a representative sequence (12.54 W mean, 25.08 W peak)**. The number matters less as a result than as an anchor: “PiM saves energy” claims in the follow-on study will be expressed as deltas against this measured joule count, closing the loop opened by the thesis’s energy motivation. For robot power budgets, the mean draw itself is notable the perception stack’s GPU share is a two-digit-watt load whose largest single component, per this analysis, is data

movement.

4.5.5 The Host-Side Boundary

Completing the end-to-end picture: per run, **66 MB is read from storage** (the sensor-data ingestion that becomes the 1.68 MB/frame H2D upload of §4.2.7), and host memory peaks at **708 MB** the keyframe/landmark store and its memory-mapped database while GPU allocation stays static at 108.65 MB (§4.4.3). The cold tier’s data path is thus measured end-to-end: storage → host memory → (fixed GPU staging buffer) → scattered on-GPU scan. Each arrow is a candidate for a different intervention (near-storage filtering, host-memory-resident scan, scatter-PiM), and the follow-on study inherits measured traffic at every stage.

4.5.6 Faults Named for the Kernels That Stay

For most kernels the verdict is “stay on GPU,” and the same evidence names each actionable fault: the loop-closure scan’s register-pressure spill (a compilation/occupancy trade, quantified at 92.6% of its DRAM volume); dependency-stalled low-occupancy kernels (G7) best addressed by launch configuration, not memory; scatter-pattern kernels whose layouts could be fixed in software before any hardware is proposed; and cache-effective kernels (G3) for which the L2 demonstrably earns its keep and must not be traded away in an NDP design. The ability to argue *against* near-data placement, kernel by kernel, is what makes the positive verdicts credible.

5 Discussion

This chapter enumerates what could be wrong, what is deliberately out of scope, and the mitigation for each. The discipline follows DAMOV’s practice of publishing limitations alongside claims.

5.1 Internal Validity

Instrumentation reactivity: Addressed directly in Chapter 4.1: profilers are trajectory-neutral (bit-identical on deterministic modes across 192 variants), the instrumented build equals the baseline build, and traces are never used for timing. Residual risk: counter collection via kernel replay perturbs caches between replays; the methodology brackets cold- and warm-cache readings rather than trusting either extreme.

Threshold choice: Classifier thresholds are frozen and sensitivity-tested at $\pm 25\%$; borderline kernels are flagged and denied high confidence. Two label-free clustering algorithms and a known-truth calibration suite independently support the class structure. Residual risk: the calibration suite’s first pass exposed three archetype-design confluences (each spanning two classes); the archetypes were corrected, classifier untouched, and the suite retained as a regression gate. The designed clock-domain intervention (§3.7) has ex ante per-class predictions and is scripted in the artifact; its execution completes the validation suite.

Attribution soundness: Three independent layers reconcile (274/274 allocations resolved); an explicit `unmapped` bucket bounds the unresolved share ($<1\text{--}7\%$ for most kernels); two systematic residuals (static module memory; texture-path reads) are named, quantified (83% of unmapped traffic in two kernels), and shown not to affect any verdict.

Nondeterminism: Long loop-closing sequences show measurable bimodality run-to-run; plain re-runs scatter identically to profiled runs, so the scatter is workload-intrinsic. All timing quantities carry a measured noise floor (CoV 0.14%).

5.2 External Validity

One workload family: cuVSLAM is a single system, argued representative by production deployment, not by breadth across SLAM implementations. Mitigations: the pipeline structure the taxonomy attaches to (streaming front-end, windowed solver, growing episodic database) is shared by keyframe-based SLAM in general; and the harness’s adapter mechanism runs arbitrary GPU workloads through the identical pipeline, so breadth is an execution cost, not a new methodology. Claims are scoped to the measured system throughout.

One primary GPU; codegen-dependent quantities: The classification survives a two-microarchitecture check at 80% agreement under a memory-system confound. However, the register-spill share (92.6% of the loop-closure scan’s DRAM volume) is a property of this compiler for this architecture: register allocation changes with target and compiler version. The thesis scopes the spill *quantity* to the measured target, and claims only the *phenomenon and method* as general, with embedded-target re-derivation (Jetson-class hardware; harness already targets it) as the first item of future work.

Dataset realism: Public benchmark corpora underrepresent degenerate lighting, dynamic scenes, and multi-camera rigs. The vendor’s multi-camera datasets are partly proprietary; the configuration regime covers modes, not every deployment condition.

5.3 Construct Validity

“PiM affinity” is candidacy, not benefit: The 61–72% offload-eligible time and the 16/49–26/49 placement counts are bounds from a first-order model with scenario parameters; no simulated or silicon NDP result is claimed. The deltas-only rule exists to keep this line uncrossed until the follow-on study.

Metric proxies: Counter metrics are proxies with known artifacts; this thesis documents a case where a proxy was right and the naive ground-truth reading was wrong (§4.3.3). The defense against proxy error is redundancy: every headline claim rests on at least two instruments that agree after reconciliation.

Texture-path blindness: The address tracer does not observe texture-unit fetches; image-tier traffic is a stated lower bound, corroborated by counter-side totals. No verdict depends on the unobserved portion.

5.4 What Is Out of Scope, and Why

Simulation of the proposed substrates (the architecture study; its configurations are generated here, its baselines measured); a second SLAM system (scope discipline: breadth via adapters instead); multi-tenant GPU sharing (cuVSLAM is single-tenant); CPU-side DAMOV replication (the origin methodology, not the contribution); and energy attribution below whole-run granularity (requires the simulator’s power model; whole-run joule count is the anchor).

6 Conclusions

6.1 What Was Established

This thesis set out to determine, by measurement, where a production visual-SLAM workload’s computation belongs on the emerging spectrum of memory-centric hardware. Its answers, in the order of the research questions:

RQ1: the methodology exists. GPU-DAMOV re-derives every component of the reference data-movement methodology for GPUs screening, machine-independent locality, threshold classification, and the full robustness battery with three GPU-driven redesigns (two-level LFMR, occupancy-conditioned latency classes, a coalescing axis) and simulated interventions replaced by measured ones. The parity argument is componentwise and complete at DAMOV’s granularity (§3.7).

RQ2: the classes are real and stable. Eight bottleneck classes (G0–G7, one dependency-bound G7 forced into existence by the data) classify 49 kernels with 91% cross-sequence consistency, survive $\pm 25\%$ threshold stress, are recovered by two label-free clustering algorithms (purity ≈ 0.68 both), agree 80% across two GPU microarchitectures, and correctly classify a known-truth calibration suite after three archetype defects were fixed with the classifier untouched.

RQ3: every byte has a name. The TaggedAllocator attribution resolves all 274 allocations and gives 48/49 kernels a unanimous data-structure tag across the corpus. The two decisive findings: the GPU memory budget is *static* (the growing keyframe database lives host-side, GPU-resident state is a fixed 6.7 MB), and 92.6% of the loop-closure scan’s DRAM volume is register-spill scratch a compiler artifact no counter-only study could have separated from data traffic.

RQ4: the placement. The workload splits three ways: a cache-immune streaming front-end (near-sensor placement; no cache size helps, measured); a hot-persistent solver that should stay on the GPU at this scale; and a cold-persistent, session-growing database tier whose scan is a genuine scattered gather (in/near-storage plus scatter-capable PiM candidacy). Time-weighted, 61–72% of GPU time carries near-data affinity; 16–26 of 49 kernels clear an analytical offload bar under conservative-to-moderate scenarios; and the baselines a follow-on study must beat

are now measured: 34.67 J whole-run energy, 1.68 MB/frame sensor upload, 41% of kernel time in boundary copies, 66 MB storage ingestion, 708 MB peak host memory.

RQ5: the measurements are trustworthy. The characterized system reproduces its vendor’s published accuracy (141-configuration matrix); profiling is output-neutral (bit-identical trajectories on deterministic modes across 192 variants); the noise floor is measured (CoV 0.14%); and the project’s two self-corrections most notably the reversal of an interim "coalesced streaming" conclusion once register-spill traffic was separated from data are documented in place, ending with independent instruments in agreement.

6.2 The Thesis, Restated

For Physical-AI perception workloads, the question "should we build PiM?" is ill-posed; the well-posed question is "*which data structure*, touched by *which kernel*, under *which access pattern*, belongs on *which substrate*?" and it is answerable today, on production software, with commodity instruments, to the granularity of a named allocation and with quantified confidence. Answered for cuVSLAM, it yields not one accelerator but a placement: sensor-side consumption for streams, the GPU for the solver, near-storage and scatter-capable near-memory engines for the map with measured evidence attached to every cell.

7 Recommendations and Future Work

7.1 Future Work

The project’s phased roadmap continues beyond this document:

1. **Complete the validation suite.** Execute the scripted clock-domain intervention (per-class predictions are stated *ex ante*: bandwidth classes track the memory clock, compute/on-chip/dependency classes the core clock) and incorporate its verdicts into the artifact.
2. **Name the last bytes.** The kernel-argument correlation layer (Layer 3) resolves the two named residual kernels carrying 83% of unmapped traffic; its target list and join interface are already committed.
3. **Embedded-target re-derivation.** Re-run the regime on Jetson-class hardware (the harness already drives remote targets) to re-derive the codegen-dependent quantities above all the spill share and test the taxonomy’s transfer to deployment-class hardware.
4. **The architecture study.** Feed the traces to Accel-Sim under the generated NDP configurations [31], with AccelWattch [32] for per-class energy, reporting speedup and energy *deltas* against this thesis’s measured baselines per the standing rule, never absolutes. The conservative/moderate placement scenarios of §4.5.3 become simulated designs; the 16–26-kernel offload sets are the experiment list.
5. **Breadth via adapters.** Run non-SLAM GPU workloads (DNN inference, GPU database operators) through the unchanged pipeline to test the methodology’s generality claim the adapter interface ensures that breadth costs only execution time, not engineering.

The characterization is complete; the design space it opens is measured, bounded, and for the first time for a production Physical-AI perception stack named, byte by byte.

8 References

- [1] G. F. Oliveira et al., ‘DAMOV: A new methodology and benchmark suite for evaluating data movement bottlenecks,’ *IEEE Access*, vol. 9, pp. 134 457–134 502, 2021.
- [2] NVIDIA Corporation, *cuVSLAM: CUDA-accelerated visual odometry and mapping*, arXiv:2506.04359, 2025.
- [3] W. A. Wulf and S. A. McKee, ‘Hitting the memory wall: Implications of the obvious,’ *ACM SIGARCH Computer Architecture News*, vol. 23, no. 1, pp. 20–24, 1995.
- [4] O. Mutlu, S. Ghose, J. Gómez-Luna and R. Ausavarungnirun, ‘Processing data where it makes sense: Enabling in-memory computation,’ *Microprocessors and Microsystems*, vol. 67, pp. 28–41, 2019.
- [5] S. Ghose, A. Boroumand, J. S. Kim, J. Gómez-Luna and O. Mutlu, ‘Processing-in-memory: A workload-driven perspective,’ *IBM Journal of Research and Development*, vol. 63, no. 6, 3:1–3:19, 2019.
- [6] J. Do, Y.-S. Kee, J. M. Patel, C. Park, K. Park and D. J. DeWitt, ‘Query processing on smart SSDs: Opportunities and challenges,’ in *ACM SIGMOD International Conference on Management of Data*, 2013, pp. 1221–1230.
- [7] B. Gu et al., ‘Biscuit: A framework for near-data processing of big data workloads,’ in *International Symposium on Computer Architecture (ISCA)*, 2016, pp. 153–165.
- [8] J. Gómez-Luna, I. El Hajj, I. Fernandez, C. Giannoula, G. F. Oliveira and O. Mutlu, ‘Benchmarking a new paradigm: Experimental analysis and characterization of a real processing-in-memory system,’ *IEEE Access*, vol. 10, pp. 52 565–52 608, 2022.
- [9] Y.-C. Kwon et al., ‘A 20nm 6GB function-in-memory DRAM, based on HBM2 with a 1.2TFLOPS programmable computing unit using bank-level parallelism, for machine learning applications,’ in *IEEE International Solid-State Circuits Conference (ISSCC)*, 2021, pp. 350–352.
- [10] S. Lee et al., ‘A 1ynm 1.25V 8Gb, 16Gb/s/pin GDDR6-based accelerator- in-memory supporting 1TFLOPS MAC operation and various activation func-

- tions for deep-learning applications,’ in *IEEE International Solid-State Circuits Conference (ISSCC)*, 2022, pp. 1–3.
- [11] S. Williams, A. Waterman and D. Patterson, ‘Roofline: An insightful visual performance model for multicore architectures,’ *Communications of the ACM*, vol. 52, no. 4, pp. 65–76, 2009.
 - [12] C. Cadena et al., ‘Past, present, and future of simultaneous localization and mapping: Toward the robust-perception age,’ *IEEE Transactions on Robotics*, vol. 32, no. 6, pp. 1309–1332, 2016.
 - [13] H. Durrant-Whyte and T. Bailey, ‘Simultaneous localization and mapping: Part I,’ *IEEE Robotics & Automation Magazine*, vol. 13, no. 2, pp. 99–110, 2006.
 - [14] J. Sturm, N. Engelhard, F. Endres, W. Burgard and D. Cremers, ‘A benchmark for the evaluation of RGB-D SLAM systems,’ in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2012, pp. 573–580.
 - [15] Z. Zhang and D. Scaramuzza, ‘A tutorial on quantitative trajectory evaluation for visual(-inertial) odometry,’ in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2018, pp. 7244–7251.
 - [16] S. Umeyama, ‘Least-squares estimation of transformation parameters between two point patterns,’ *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 13, no. 4, pp. 376–380, 1991.
 - [17] J. Ahn, S. Hong, S. Yoo, O. Mutlu and K. Choi, ‘A scalable processing-in-memory accelerator for parallel graph processing,’ in *International Symposium on Computer Architecture (ISCA)*, 2015, pp. 105–117.
 - [18] J. Ahn, S. Yoo, O. Mutlu and K. Choi, ‘PIM-enabled instructions: A low-overhead, locality-aware processing-in-memory architecture,’ in *International Symposium on Computer Architecture (ISCA)*, 2015, pp. 336–348.
 - [19] O. Villa, M. Stephenson, D. Nellans and S. W. Keckler, ‘NVBit: A dynamic binary instrumentation framework for NVIDIA GPUs,’ in *IEEE/ACM International Symposium on Microarchitecture (MICRO)*, 2019, pp. 372–383.

- [20] J. Cao, R. Sarkar, R. Kumar, J. Arulraj and H. Kim, *GPU database systems characterization and optimization*, Proceedings of the VLDB Endowment, 2023.
- [21] R. L. Mattson, J. Gecsei, D. R. Slutz and I. L. Traiger, ‘Evaluation techniques for storage hierarchies,’ *IBM Systems Journal*, vol. 9, no. 2, pp. 78–117, 1970.
- [22] K. Hsieh et al., ‘Transparent offloading and mapping (TOM): Enabling programmer-transparent near-data processing in GPU systems,’ in *International Symposium on Computer Architecture (ISCA)*, 2016, pp. 204–216.
- [23] A. Boroumand et al., ‘Google workloads for consumer devices: Mitigating data movement bottlenecks,’ in *International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, 2018, pp. 316–331.
- [24] R. Mur-Artal and J. D. Tardós, ‘ORB-SLAM2: An open-source SLAM system for monocular, stereo, and RGB-D cameras,’ *IEEE Transactions on Robotics*, vol. 33, no. 5, pp. 1255–1262, 2017.
- [25] C. Campos, R. Elvira, J. J. G. Rodríguez, J. M. M. Montiel and J. D. Tardós, ‘ORB-SLAM3: An accurate open-source library for visual, visual-inertial, and multimap SLAM,’ *IEEE Transactions on Robotics*, vol. 37, no. 6, pp. 1874–1890, 2021.
- [26] L. Nardi et al., ‘Introducing SLAMBench, a performance and accuracy benchmarking methodology for SLAM,’ in *IEEE International Conference on Robotics and Automation (ICRA)*, 2015, pp. 5783–5790.
- [27] A. Geiger, P. Lenz, C. Stiller and R. Urtasun, ‘Vision meets robotics: The KITTI dataset,’ *International Journal of Robotics Research*, vol. 32, no. 11, pp. 1231–1237, 2013.
- [28] M. Burri et al., ‘The EuRoC micro aerial vehicle datasets,’ *International Journal of Robotics Research*, vol. 35, no. 10, pp. 1157–1163, 2016.
- [29] D. Schubert, T. Goll, N. Demmel, V. Usenko, J. Stückler and D. Cremers, ‘The TUM VI benchmark for evaluating visual-inertial odometry,’ in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2018, pp. 1680–1687.

- [30] A. Handa, T. Whelan, J. McDonald and A. J. Davison, ‘A benchmark for RGB-D visual odometry, 3D reconstruction and SLAM,’ in *IEEE International Conference on Robotics and Automation (ICRA)*, 2014, pp. 1524–1531.
- [31] M. Khairy, Z. Shen, T. M. Aamodt and T. G. Rogers, ‘Accel-sim: An extensible simulation framework for validated GPU modeling,’ in *International Symposium on Computer Architecture (ISCA)*, 2020, pp. 473–486.
- [32] V. Kandiah et al., ‘AccelWattch: A power modeling framework for modern GPUs,’ in *IEEE/ACM International Symposium on Microarchitecture (MICRO)*, 2021, pp. 738–753.
- [33] J. H. Ward, ‘Hierarchical grouping to optimize an objective function,’ *Journal of the American Statistical Association*, vol. 58, no. 301, pp. 236–244, 1963.

9 Appendices

9.1 Glossary

Arithmetic intensity (AI) FLOPs performed per byte fetched from DRAM; the x-axis of the roofline. Low AI = memory-bound.

Attribution Assigning each memory access to the named data structure (allocation) it falls within.

Bandwidth Data volume movable per unit time (GB/s); contrast *latency*.

Bank (DRAM) A subdivision of a DRAM device; near-bank PiM places compute at each bank to harvest internal bandwidth.

Bundle adjustment (BA) Joint nonlinear least-squares refinement of camera poses and 3-D landmarks; cuVSLAM’s numerical core.

Cache (L1/L2) Automatic on-chip copies of recently used data; L1 per SM, L2 chip-wide (25 MB on the primary instrument).

Coalescing Hardware merging of a warp’s 32 lane addresses into few 32-byte sectors; its failure mode (scatter) wastes up to $8\times$ bandwidth.

Cold-persistent Data with session lifetime touched episodically (the keyframe database).

Compulsory miss A first-touch cache miss no cache size can avoid; streaming data consists almost entirely of these.

CoV Coefficient of variation (standard deviation $\div$ mean); run-to-run noise measure (0.14% at locked clocks here).

DRAM speed-of-light (DRAM SoL) Achieved DRAM bandwidth as a percentage of the measured peak; $\geq 50\%$ is treated as saturation.

Footprint The number of distinct 32-byte sectors a launch touches, i.e. its exact working-set size.

G0–G7 The GPU bottleneck taxonomy (Table 3).

Ground truth The externally measured true trajectory against which SLAM output is scored.

Hot-persistent Resident data reused every frame (the BA linear system).

Jaccard overlap $|A \cap B| / |A \cup B|$ between consecutive launches' footprints; measures working-set migration.

Kernel One GPU function launched over a grid of threads.

Keyframe A frame promoted into the long-term map; its descriptors enter the loop-closure database.

Lane One of the 32 SIMT execution slots of a warp.

LFMR Last-to-first miss ratio; on a two-level GPU hierarchy, $1 - \text{L2 hit rate}$. Near 1 = caches not helping.

Local memory Per-thread DRAM-backed space used for register spills; coalesced by construction; *not* program data.

Loop closure Recognition of a previously visited place, triggering trajectory-wide drift correction via a database scan.

Memory wall The compounding gap between compute throughput and memory speed; the thesis's founding observation.

MPKI DRAM sector fetches per thousand warp instructions.

Occupancy Resident warps as a fraction of the SM maximum; the GPU's latency-hiding budget.

Persistence The lifetime/touch-pattern axis of a data structure: streaming, hot-persistent, or cold-persistent.

PiM / NDP Processing-in-memory / near-data processing; compute placed at or near DRAM.

Reuse distance Distinct sectors touched between successive accesses to the same sector; its CDF predicts hit rate for any cache size.

Roofline $\min(\text{peak compute}, \text{AI} \times \text{peak bandwidth})$ performance bound; placement under the sloped segment = memory-bound.

Sector The 32-byte granule of GPU DRAM/L2 access; four per 128-byte line.

Sectors per request Average sectors fetched per warp memory request; ~ 4 coalesced, $\rightarrow 32$ scattered.

SM Streaming multiprocessor; the GPU’s core unit (22 on the primary instrument).

Spill (register) Compiler-generated traffic to local memory when per-thread state exceeds the register budget.

Stall reason Hardware-recorded cause of a warp’s failure to issue (memory scoreboard, dependency wait, barrier, ...).

Streaming Per-frame data consumed once; compulsory-miss dominated.

Substrate A candidate execution home: GPU, CPU, DRAM-PiM, scatter-PiM, near-sensor, in/near-storage.

TaggedAllocator This project’s allocation-journaling instrumentation: pointer, size, call stack, and data-structure tag for every GPU allocation.

Warp The fixed 32-thread SIMT execution bundle.

9.2 Artifact and Reproducibility Description

All software, configurations, and derived data of this thesis are archived in a single repository (github.com/mij001/cuvslam-stack, MIT license). The artifact is organized so that *every table in this document regenerates from committed CSV data using only the Python standard library no GPU, no datasets, no vendor tools required*; capture (which does require the hardware) is separated from analysis by construction.

9.2.1 Layout

configs/ The configuration regime: a small set of human-owned base TOML files plus a deterministic mutation pass that materializes the full experiment matrices (141 accuracy configurations, 192 profiling-coverage variants). One TOML fully describes a run: dataset, camera rig, pipeline mode, feature toggles, and the ground-truth evaluation block.

profiling/harness/ The capture entry point: one command runs any configuration under `nsys`, `ncu` (curated metric set, bounded windows), or `NVBit` (windowed, kernel-filtered address traces), writing versioned result directories with full provenance (clock locks, driver/tool versions, command lines). A workload-*adapter* interface makes the harness generic: `cuVSLAM` is one adapter; a command adapter runs arbitrary GPU programs through the identical pipeline.

profiling/analysis/ The stdlib-only analysis library: screening, classification (the decision tree of §3.4, with thresholds defined once and imported live into documentation), roofline, locality (reuse-distance CDFs with memory-space filtering), attribution (the journal-trace join), campaign synthesis, substrate verdicts, residual ranking, and the DAMOV-parity checks. The GPU-free test suite fabricates known inputs for every module.

profiling/calibration/ The known-truth archetype suite (§4.2.5): eight kernels, one per class, with the blind-classification driver.

profiling/validation/ The clock-domain intervention scripts (§3.7) and the profiler-neutrality drivers.

patches/ The instrumentation as reviewable diffs: the `TaggedAllocator` + `NVTX` enablement patch against the `cuVSLAM` source tree, and the `NVBit` `mem_trace` extensions (launch-window and kernel-name gating; the allocation-event sidecar).

reports/, profiling/reports/ Dated, immutable result directories: the 27-sequence characterization campaign, the locality study and its space-filtered re-derivation (both the superseded and corrected tables are retained, with the correction

marked in place), the attribution campaign, the 141-run accuracy matrix, the 192-variant neutrality campaign, the substrate synthesis, the energy and host-I/O measurements, and the DAMOV-validation artifacts (cross-device agreement, hierarchical-clustering agreement).

dashboard/ The evidence explorer: an interactive application presenting every finding with its computed number, the per-run reasoning chain (time share $\rightarrow$ metrics vs. thresholds $\rightarrow$ the decision-tree rule that fired $\rightarrow$ verdict), and a methodology tab whose formulas and thresholds are stamped from the analysis source at build time so interface and code cannot disagree.

9.2.2 Environment

Primary testing hardware: NVIDIA RTX 2000 Ada (`sm_89`), driver 575.64.05, CUDA 12.9.1, Nsight Compute 2025.2.1.3, Nsight Systems 2025.3.2, and NVBit 1.8. Maintaining this exact combination of software and driver versions is critical, since the binary instrumentation specifies the driver version, which, in its turn, determines the profiler versions. To ensure compatibility, an environment doctor checks the target system and suggests a one-line fix for any incompatibilities detected. GPU clocks are locked (1620/7001 MHz) for every measurement; ceilings are re-measured per instrument.

9.2.3 Reproduction paths

Analysis-only (no GPU): clone; run the test suite; regenerate any table in this thesis from the committed CSVs via the corresponding `analysis` module. **Full capture:** configure a target with the environment doctor; run the regime’s campaign drivers (all are resumable and idempotent re-running skips completed steps); the coverage audit then proves the captures complete before analysis will consume them. **Instrumented builds:** apply the committed patch to the cu-VSLAM source tree and build with the pinned container recipe; the baseline and instrumented binaries are kept side by side, and the QoR check of §4.1.2 re-verifies output equivalence on any rebuild.

9.3 Extended Result Tables

The tables below are excerpts of committed artifacts, reproduced for convenience; the full machine-readable versions live in the repository (Appendix 9.2).

9.3.1 Cross-sequence attribution consistency (excerpt)

Selected rows of the 49-kernel attribution-consistency table: each kernel’s modal top global-space tag, the agreement across every sequence it appears in, and its median memory-space split (shared/spill/global).

Table 9: Attribution consistency across 27 sequences (excerpt).

kernel	modal global tag	agree	shared	spill	global
cast_image_kernel	images_raw	100%	0%	0%	100%
cast_image_kernel_rgb	pyramid_levels	100%	0%	0%	100%
gaussian_scaling_kernel	images_raw	100%	0%	0%	100%
conv_grad_x_kernel	pyramid_levels	100%	92%	0%	8%
conv_grad_y_kernel	pyramid_levels	100%	96%	0%	4%
gfft_values_kernel	feature_tracks	100%	98%	0%	2%
non_max_suppression	feature_tracks	100%	90%	0%	10%
accumulateGFFT_kernel	feat. sel. scratch	100%	89%	0%	11%
sba::build_full_system_1	ba_linear_system	100%	0%	0%	100%
sba::calc_jacobians	ba_linear_system	100%	4%	0%	96%
sba::reduced_system_st. 2	ba_linear_system	100%	4%	0%	96%
st_build_cache	keyframe_descriptors	100%	0%	0%	100%
st_track_with_cache	keyframe_descriptors	100%	0%	93%	7%
cub::DeviceMergeSort*	feature_tracks	100%	83–97%	0%	3–17%

9.3.2 Instrumented vs. baseline trajectories (QoR excerpt)

9.3.3 KITTI per-sequence odometry accuracy

9.3.4 Measured stage map (NVTX, excerpt)

Table 10: Quality-of-results check: instrumented (TaggedAllocator) build vs. baseline build, same configurations (APE, m).

configuration	baseline	instrumented	Δ
EuRoC MH_01 inertial SLAM	0.689	0.689	0.0000
EuRoC V1_01 stereo SLAM	0.037	0.037	0.0000
EuRoC V2_02 stereo odometry	0.177	0.177	0.0000
KITTI 00 stereo odometry	6.841	6.605	-0.236^a
KITTI 06 stereo SLAM	2.312	2.277	-0.035^a
TUM long_office RGB-D SLAM	0.018	0.021	$+0.003$

^a Within the sequence’s own plain-rerun scatter (§4.1.3); kitti00 is the canonical bimodal loop-closure case.

Table 11: KITTI stereo odometry, per sequence: absolute error and segment-relative drift (100–800 m KITTI protocol).

seq.	APE (m)	avgRTE (%)	avgRE (°)
00	6.84	0.892	0.938
01	11.59	1.899	0.386
02	4.55	0.975	0.707
03	3.94	2.265	0.988
04	2.16	1.967	0.090
05	3.31	0.974	0.723
06	2.88	1.248	0.757
07	1.44	0.982	1.131
08	4.57	1.221	0.959
09	3.53	1.238	0.544
10	2.17	1.010	0.661

At the 500 m segment length the corpus-level translational drift is 0.82%, against the 0.85% published KITTI-protocol figure (§4.1.1). Sequence 01 (highway, sparse features) dominates the absolute-APE average; sequences 01 and 04 close no loops and serve as the loop-free control in the attribution campaign.

Table 12: Kernel-to-stage mapping, measured from cuVSLAM’s own NVTX ranges.

NVTX stage	kernels
front-end, per frame	cast_image_kernel_rgb, gaussian_scaling, conv_grad_x/y, lk_track
keyframe feature refresh	gftt_values, accumulateGFTT, downsample_gftt, non_max_suppression, select_features, CUB keypoint sort
visual-odometry solver	matcher::*, cast_depth, sba::reduced_system_*, cuSOLVER helpers
windowed bundle adjustment	sba::build_full_system_*, sba::calc_jacobians, sba::evaluate_cost
keyframe ingest	st_build_cache
loop closure & optimization	st_track_with_cache